

August 28

K E Y S P A C E  
Amsterdam

# Demystifying Valkey Clustering: Architecture, Fault Tolerance, and Scalability

Harkrishn Patro

Senior Software Engineer - AWS/Valkey Maintainer



# Agenda

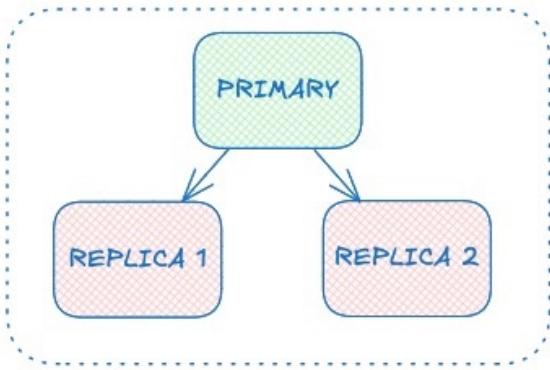
- Overview
- Operation mode
- Cluster mode components
- Cluster coordination mechanisms
- Recent improvements
- Benchmark
- Future work

# Introducing the Valkey project

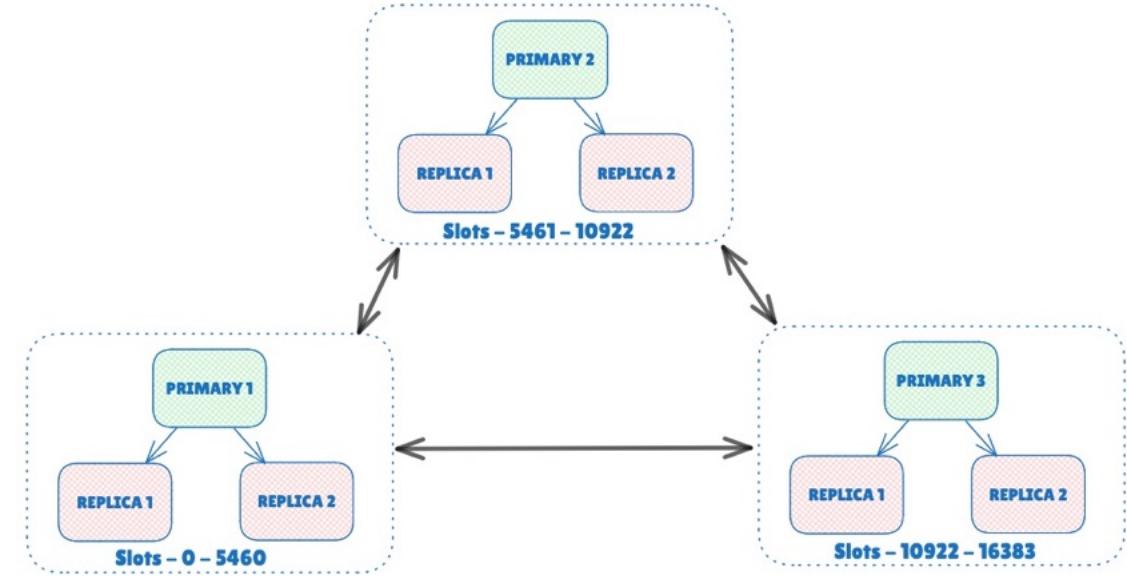
- In-memory key value store
- Supports over 200 commands and multiple data structures like hashes, sets, sorted sets
- Used for caching, session management, leaderboard, message broker



# Operation mode



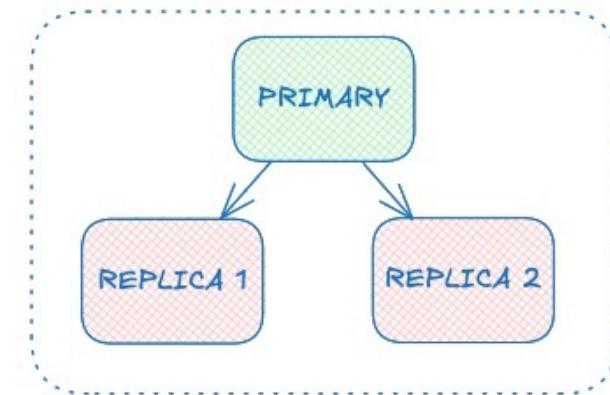
STANDALONE



CLUSTER

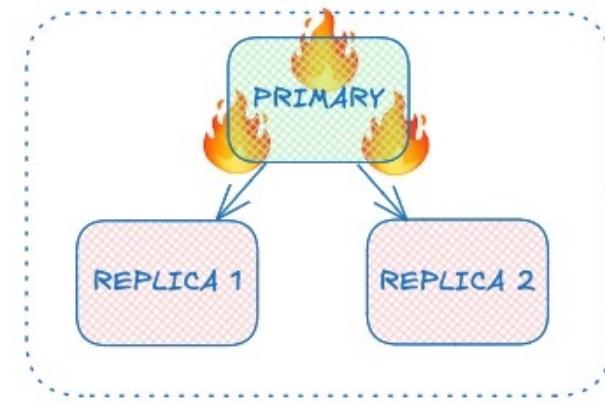
# Valkey Standalone Overview

- Single keyspace
- Asynchronous replication
- Scalable reads



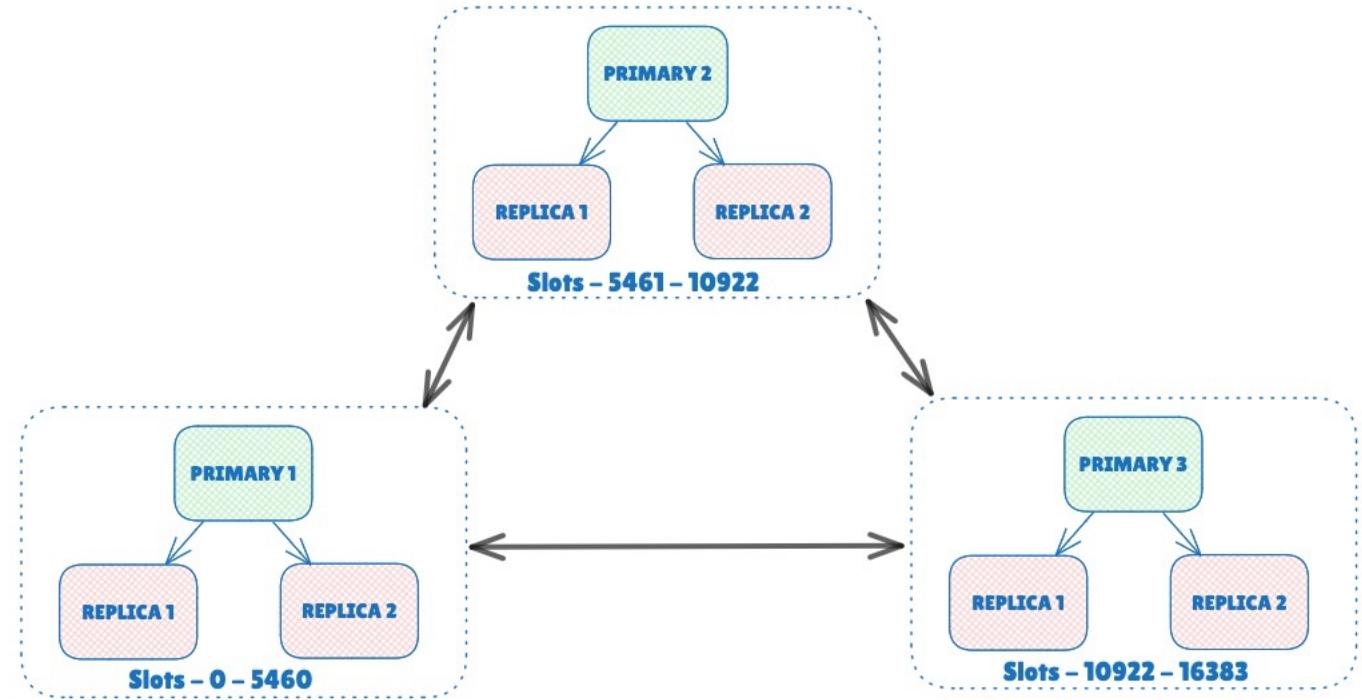
# Valkey Standalone - Summary

- Easy to setup
- Lack of fault tolerance
- Memory and write throughput bottleneck



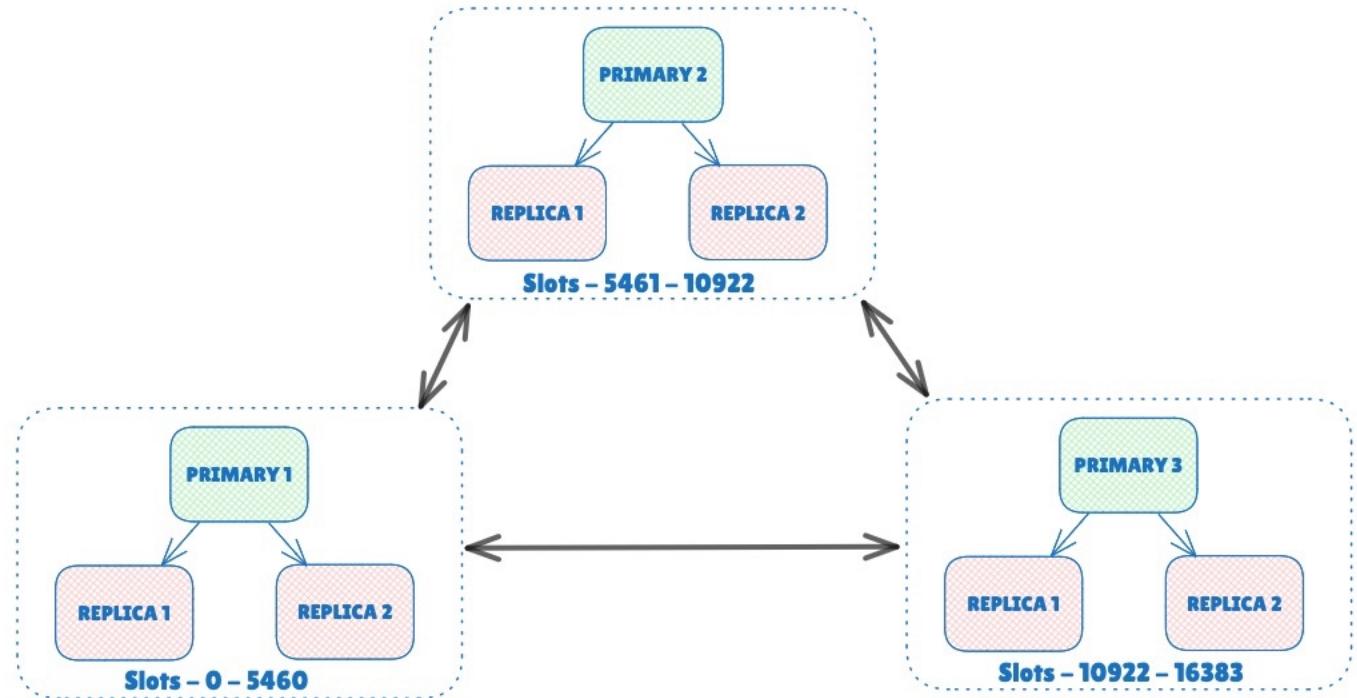
# Valkey Cluster Overview

- Horizontally scalable
- Homogenous setup
- Server enabled data sharding
- Automatic client redirection
- In-built fault tolerance



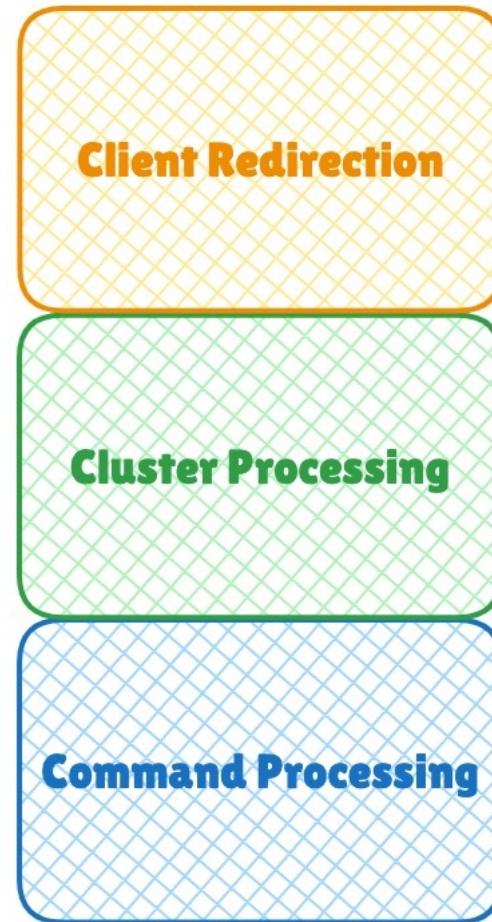
# Cluster Mode Components

- Cluster Node
- Cluster Slots
- Cluster Bus



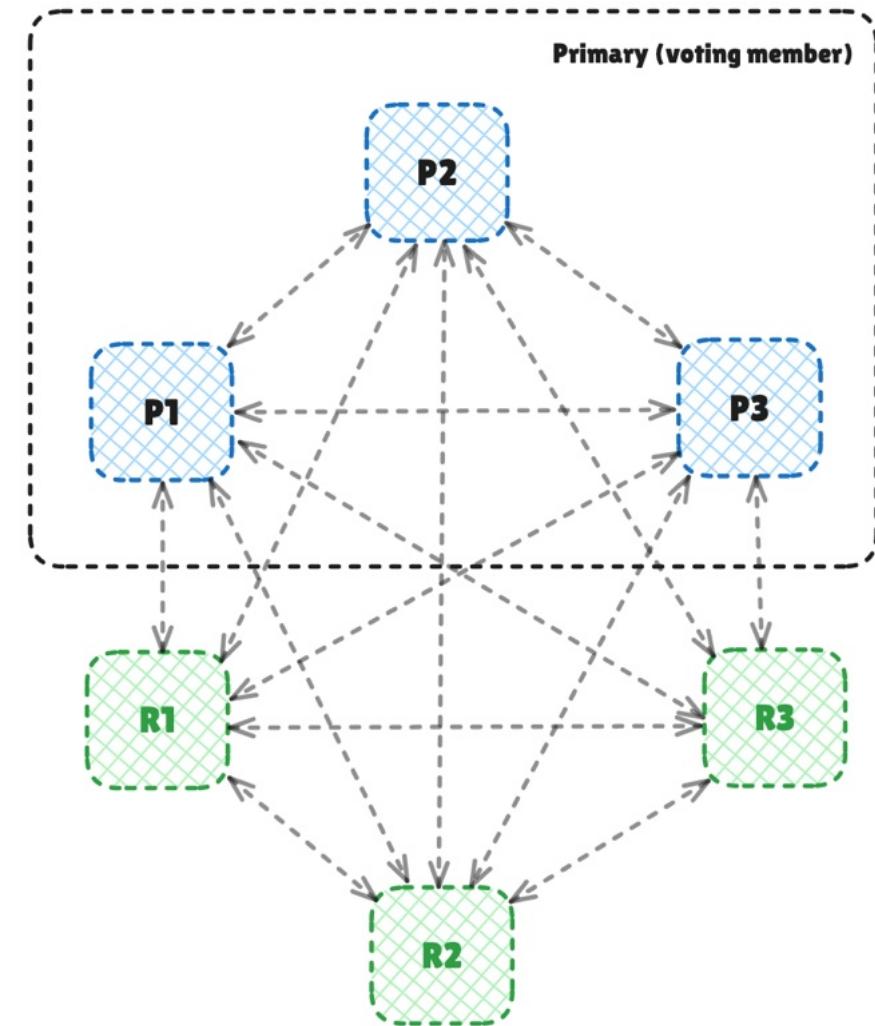
# Cluster Node

- Serves data
- Health detection
- Serves topology information



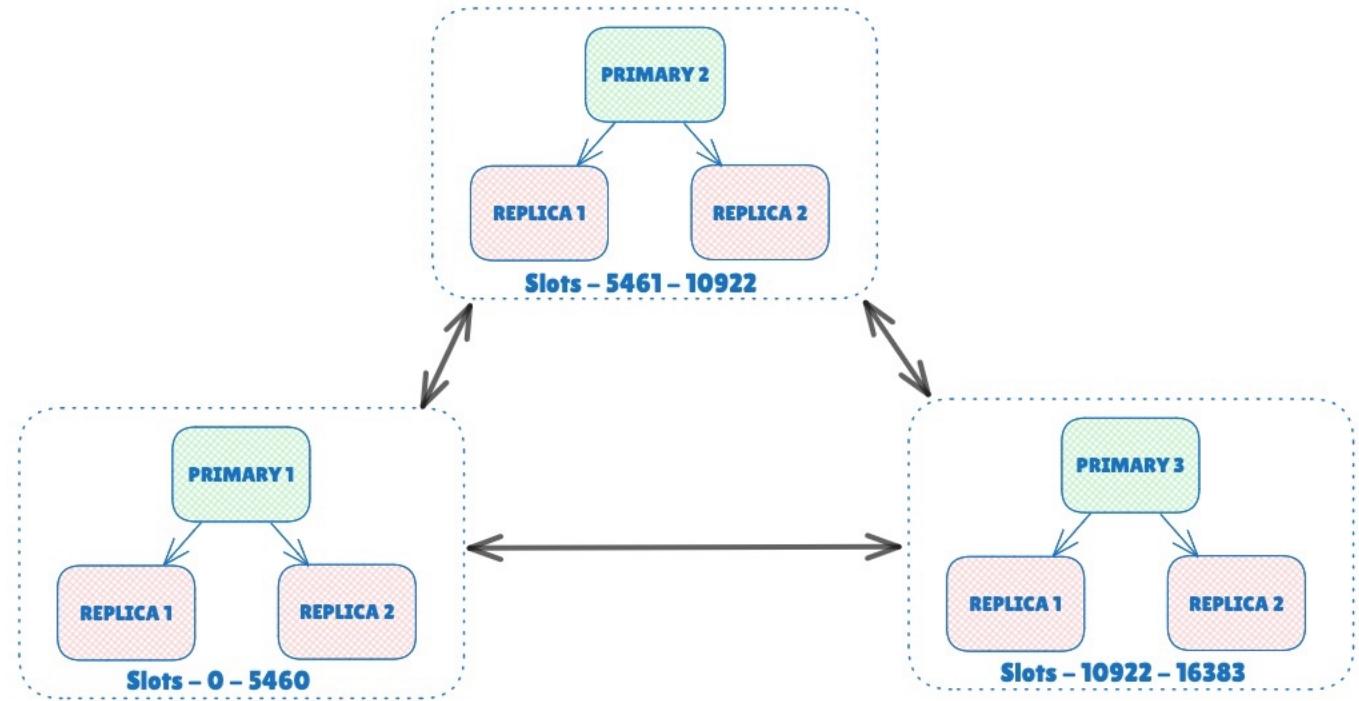
# Cluster Node (Contd.)

- Primary
  - Quorum - Voting member
- Replica
  - Election candidate



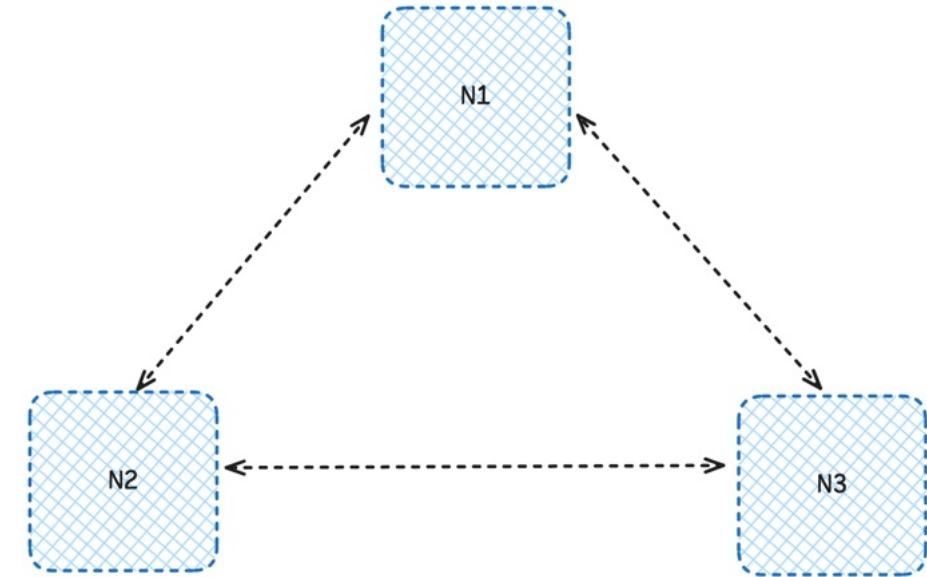
# Cluster Slots

- Hashes key to a slot
- Slot to a node
- Key distribution – 16384 slots
- Algorithm -  $\text{CRC16}(\text{key}) \% 16384$
- Primary – one or more slots



# Cluster Bus

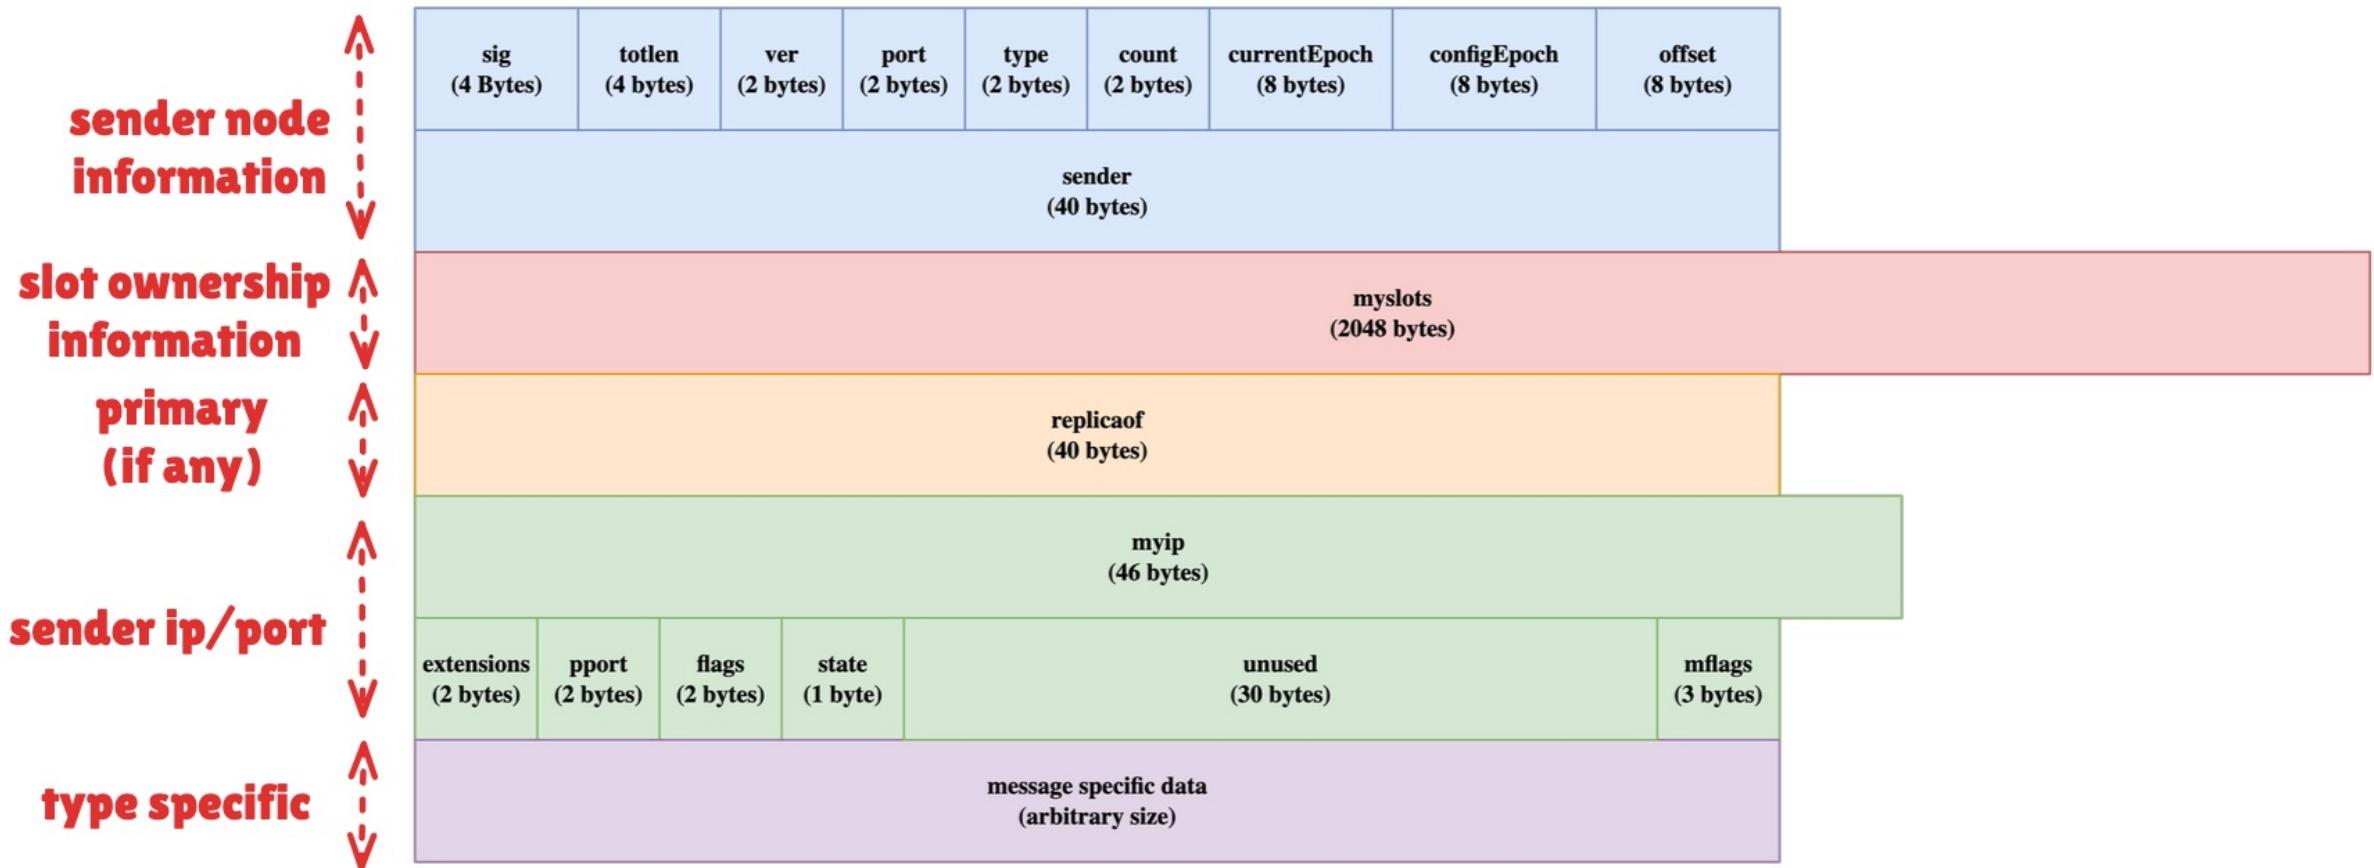
- Bidirectional persistent TCP connection between nodes
- Mesh topology
- Custom message protocol
- Gossip – piggyback information
- Supports Pub/Sub traffic



# Cluster Bus – Message Type

- New node discovery - MEET
- Heart beat – PING / PONG
- Node failure – FAIL / UPDATE
- PubSub – PUBSUB / PUBSUBSHARD

# Cluster Bus - Message Header Format



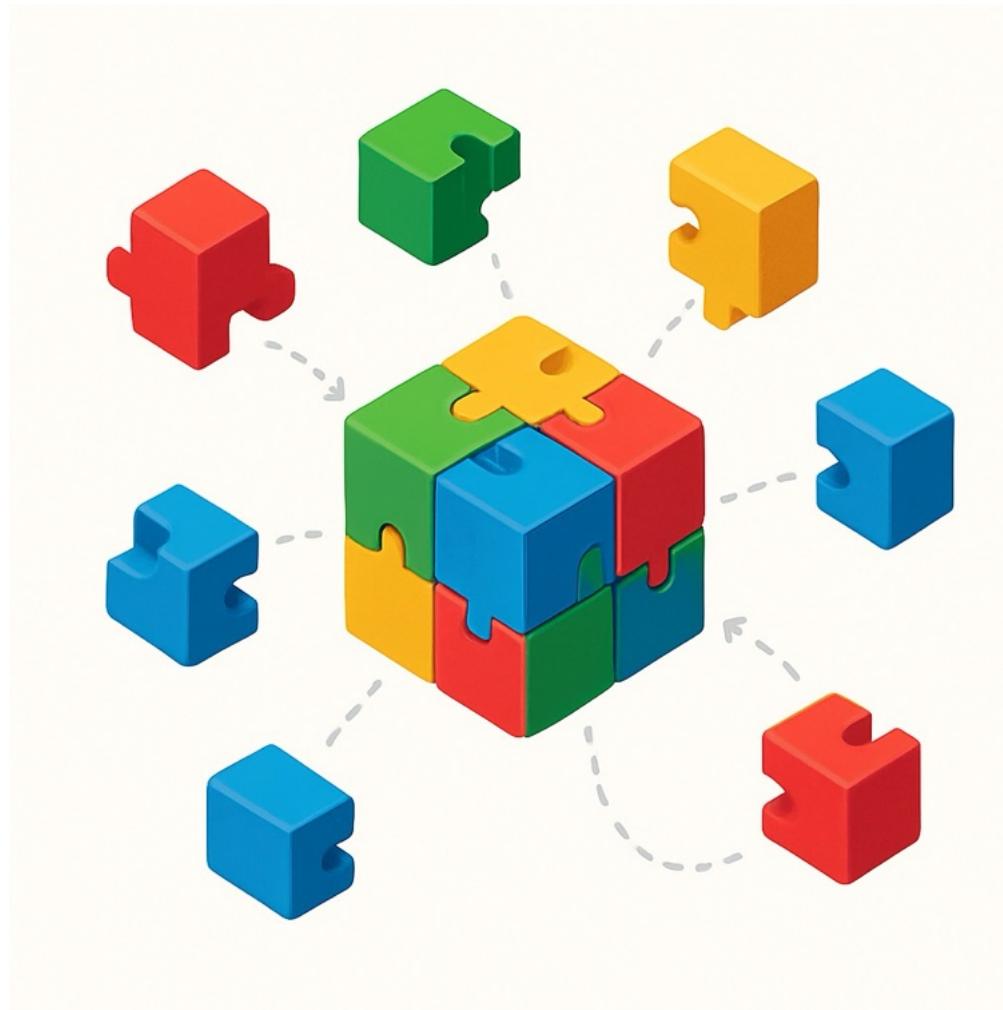
# Cluster Bus - Message Header (PING/PONG/MEET)

## Fixed Message Header (~ 2KB)

N array gossip							
nodename (40 bytes)	ping_sent (4 bytes)	pong_received (4 bytes)	ip (46 bytes)	port (2 bytes)	cport (2 bytes)	flags (2 bytes)	pport (2 bytes)
nodename	ping_sent	pong_received	ip	port	cport	flags	pport
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.
nodename (40 bytes)	ping_sent (4 bytes)	pong_received (4 bytes)	ip (46 bytes)	port (2 bytes)	cport (2 bytes)	flags (2 bytes)	pport (2 bytes)

# Cluster Coordination Mechanism

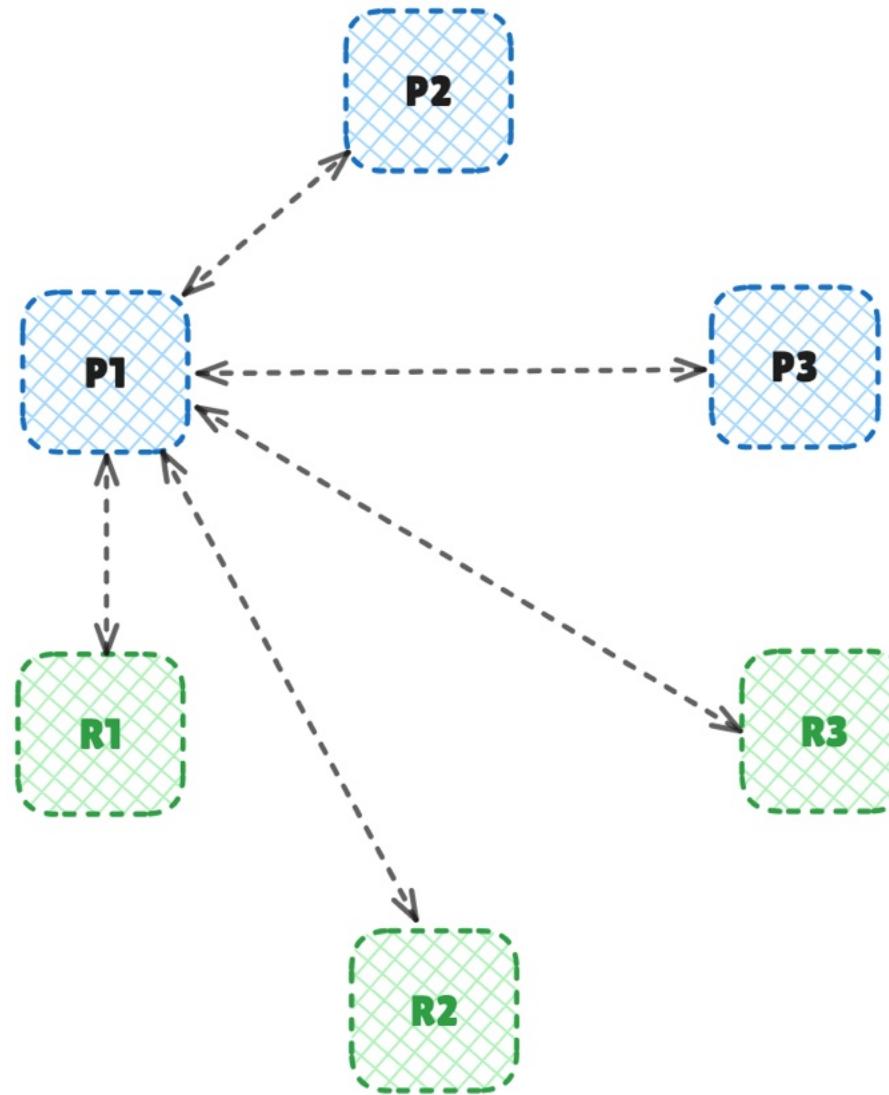
- Membership
- Redirection - Client
- Failure detection
- Failover
- Versioning – Conflict Resolution



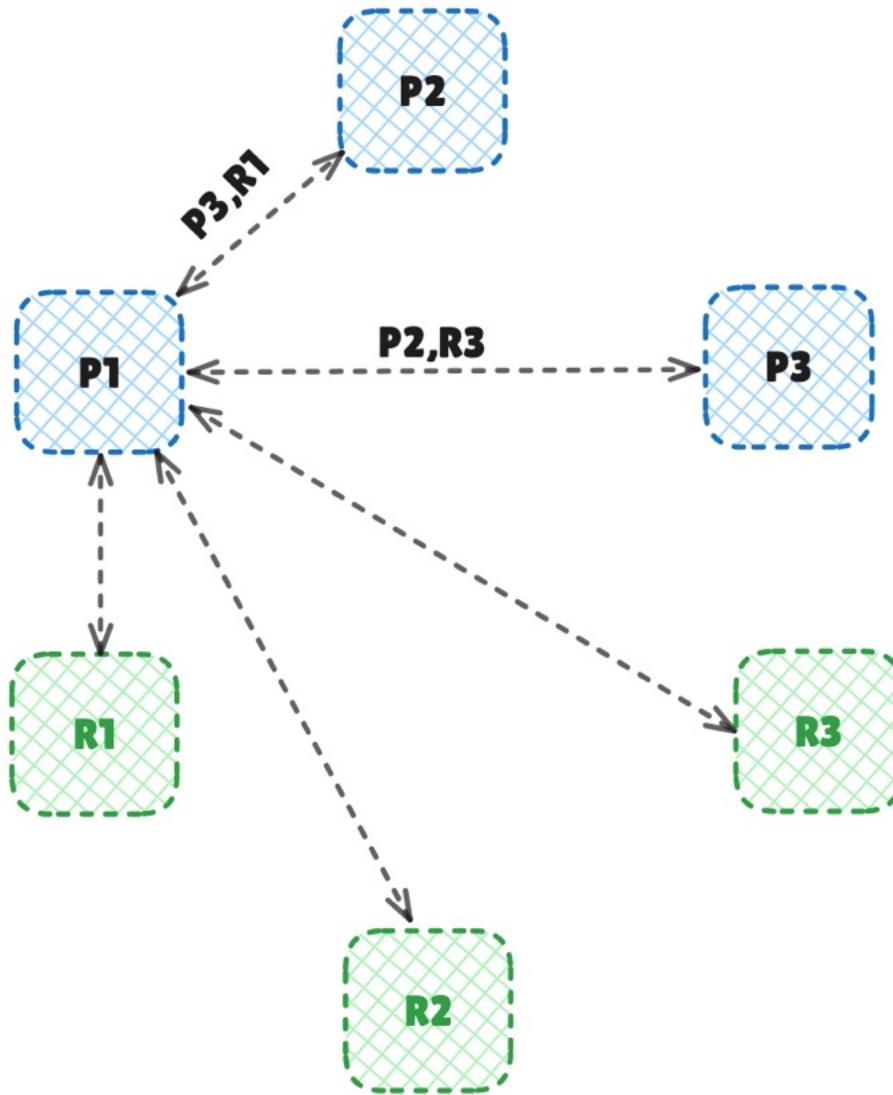
# Cluster Node Membership



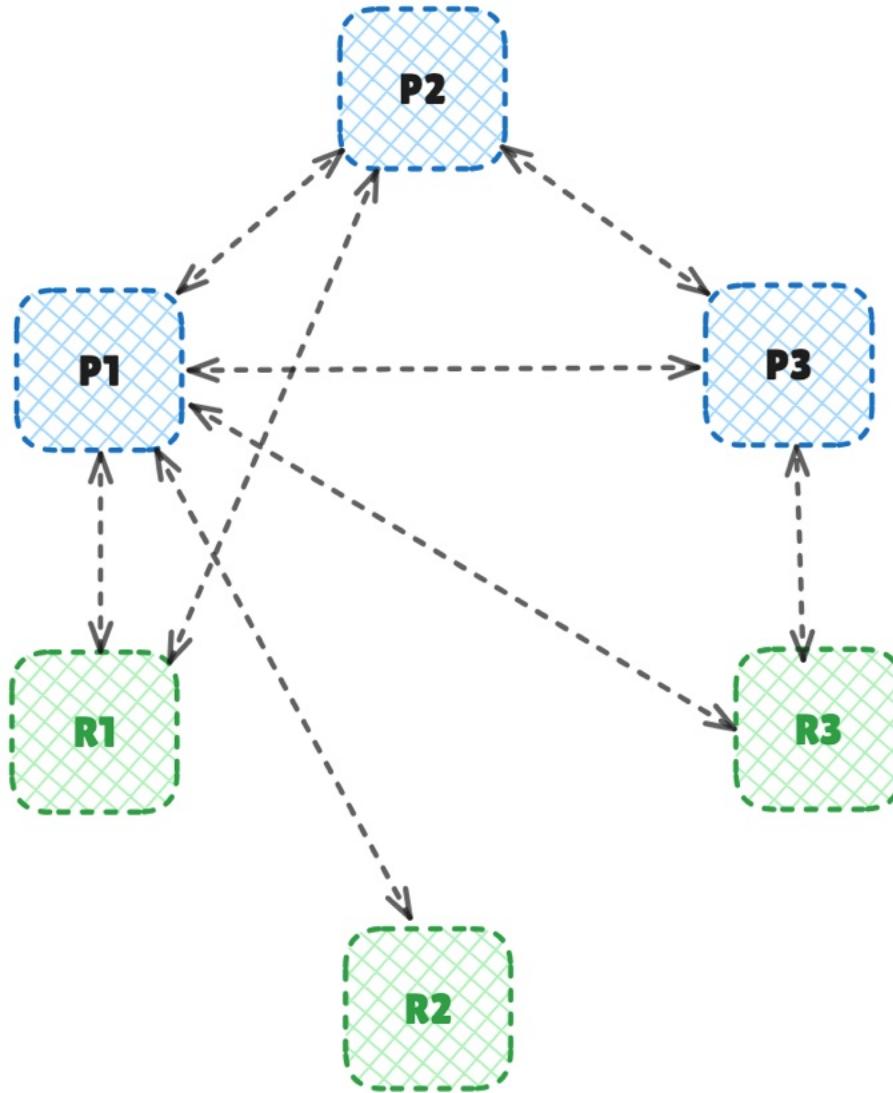
# Node Membership – Seed node discovery



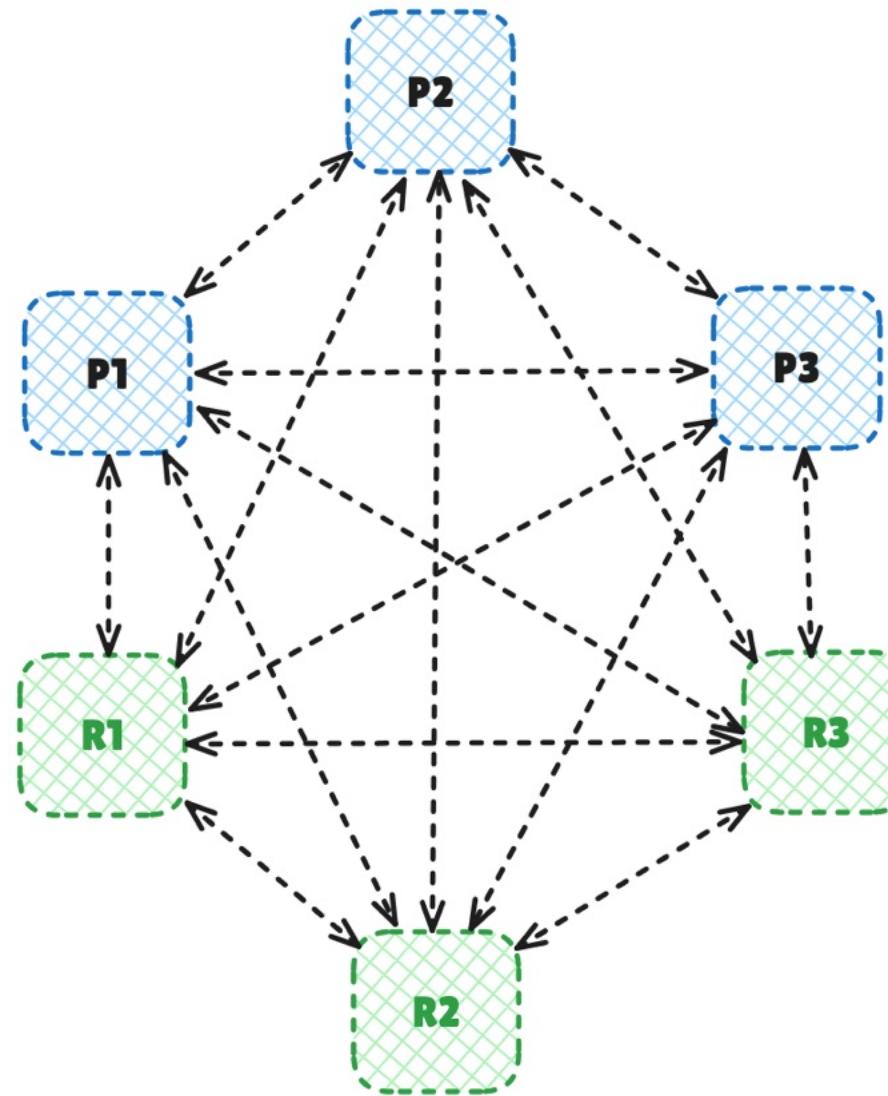
# Node Membership – Gossip



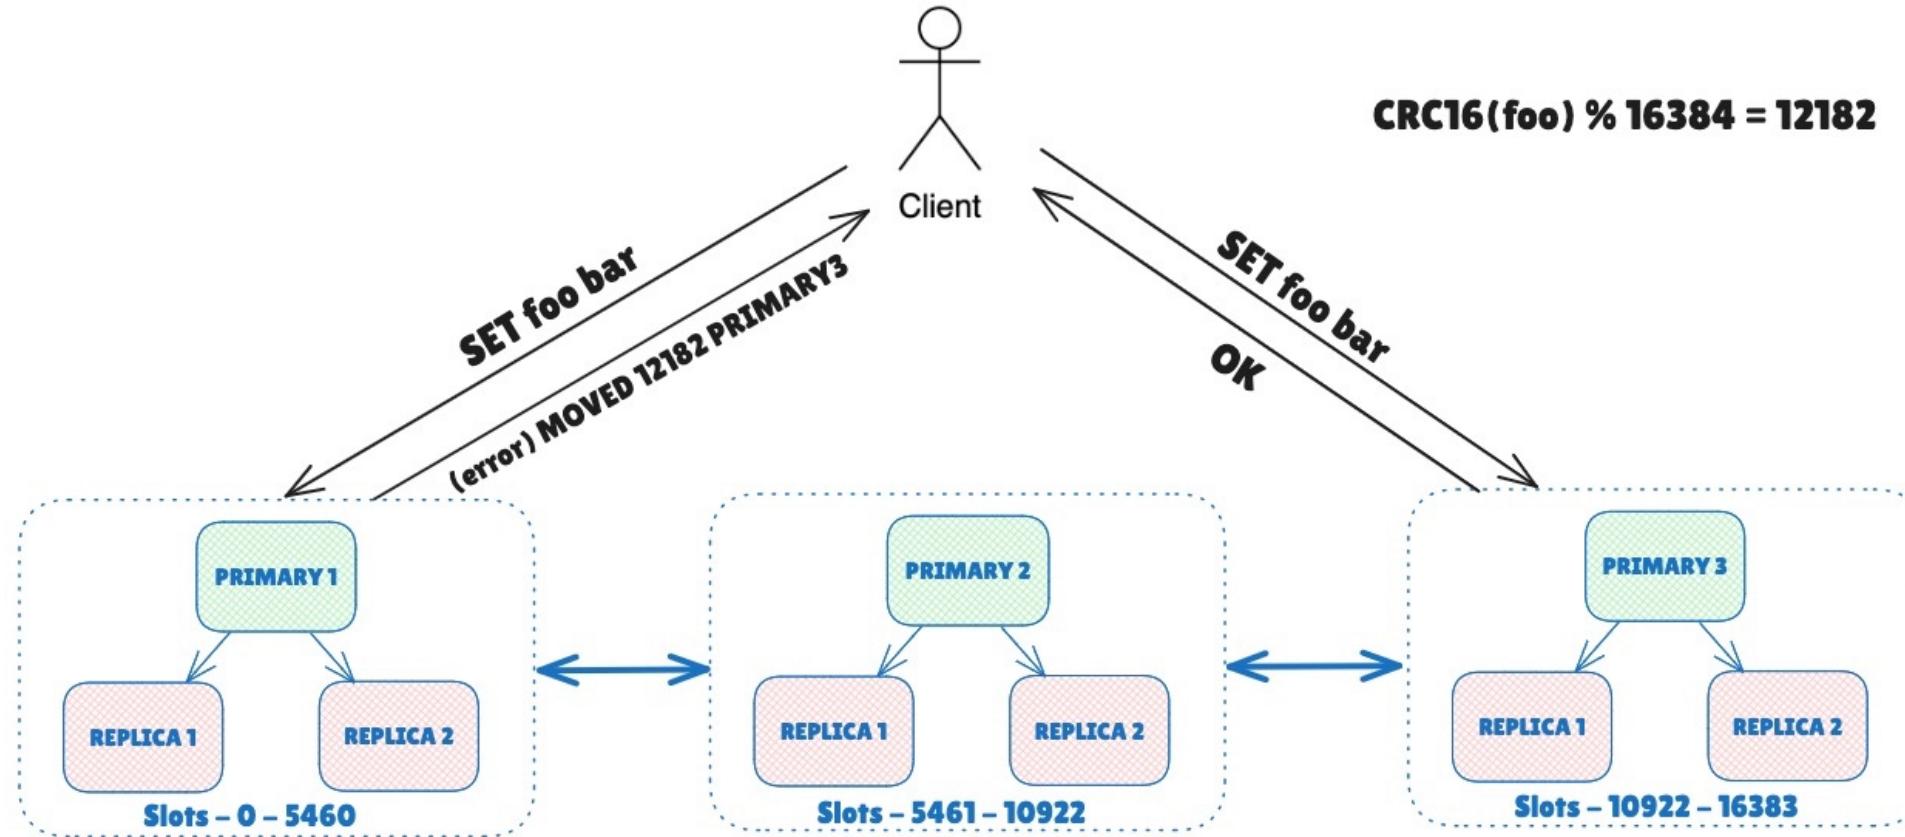
# Node Membership- Discovery via gossip



# Node Membership – Fully connected



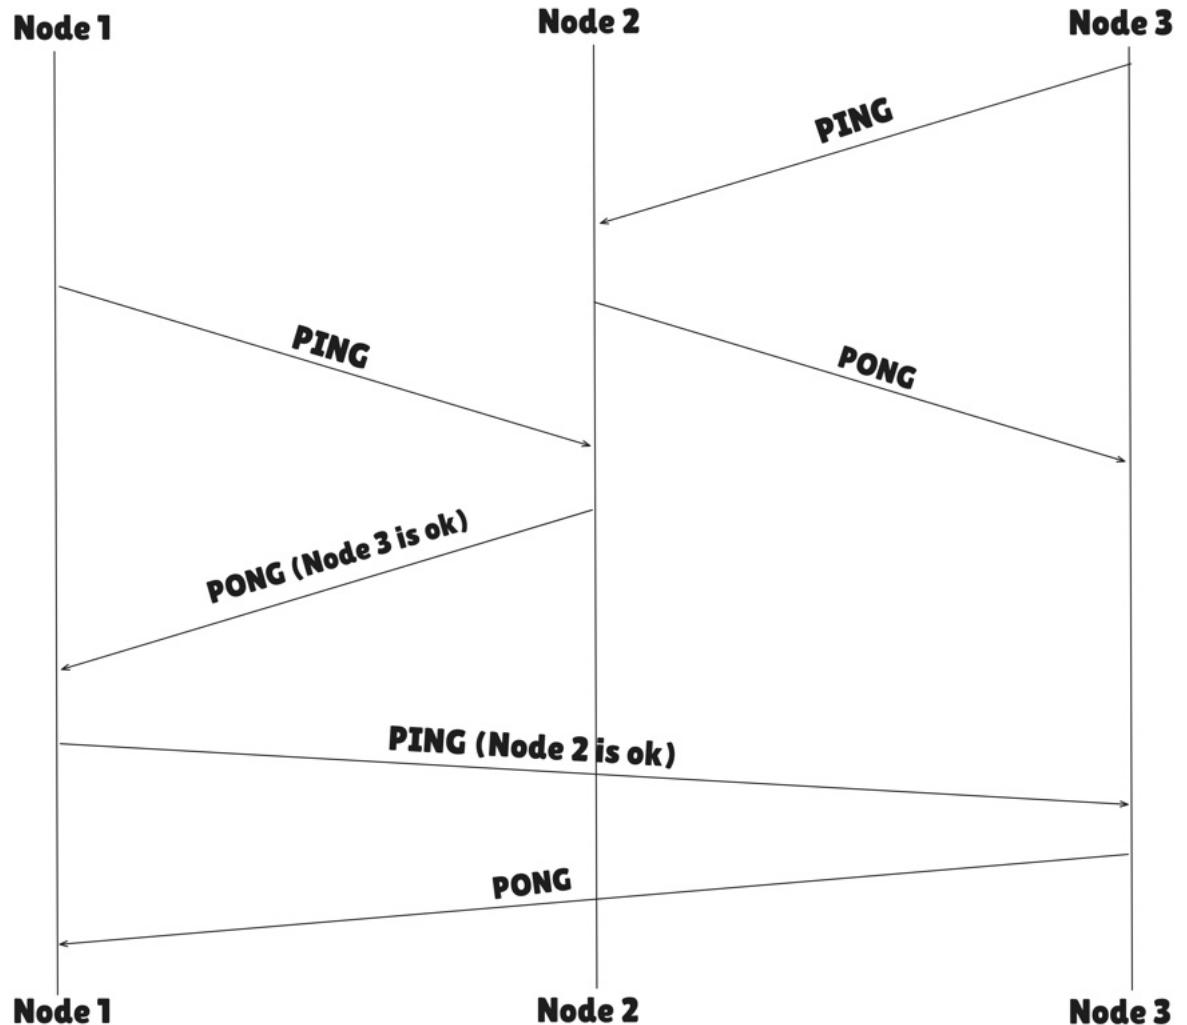
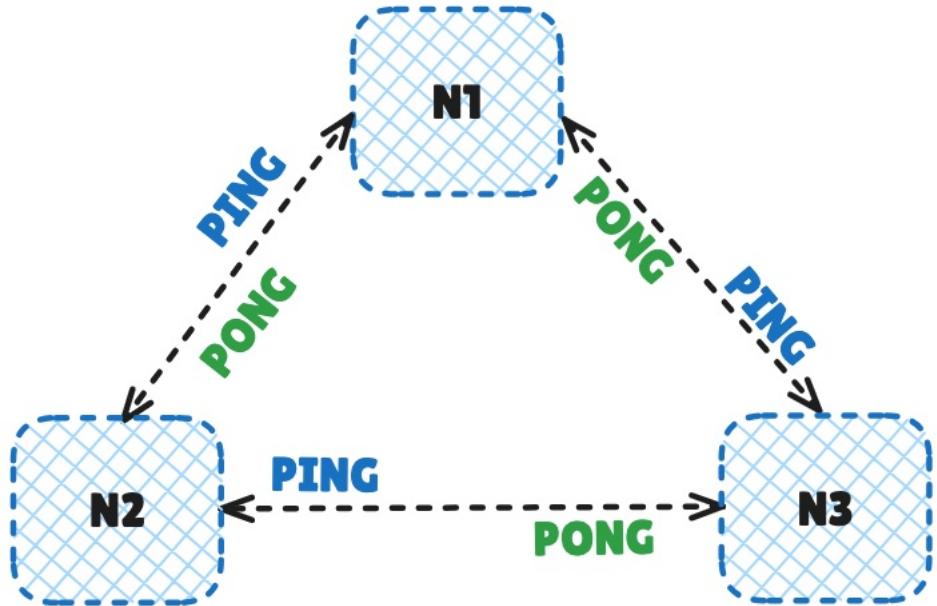
# Node Redirection - Client



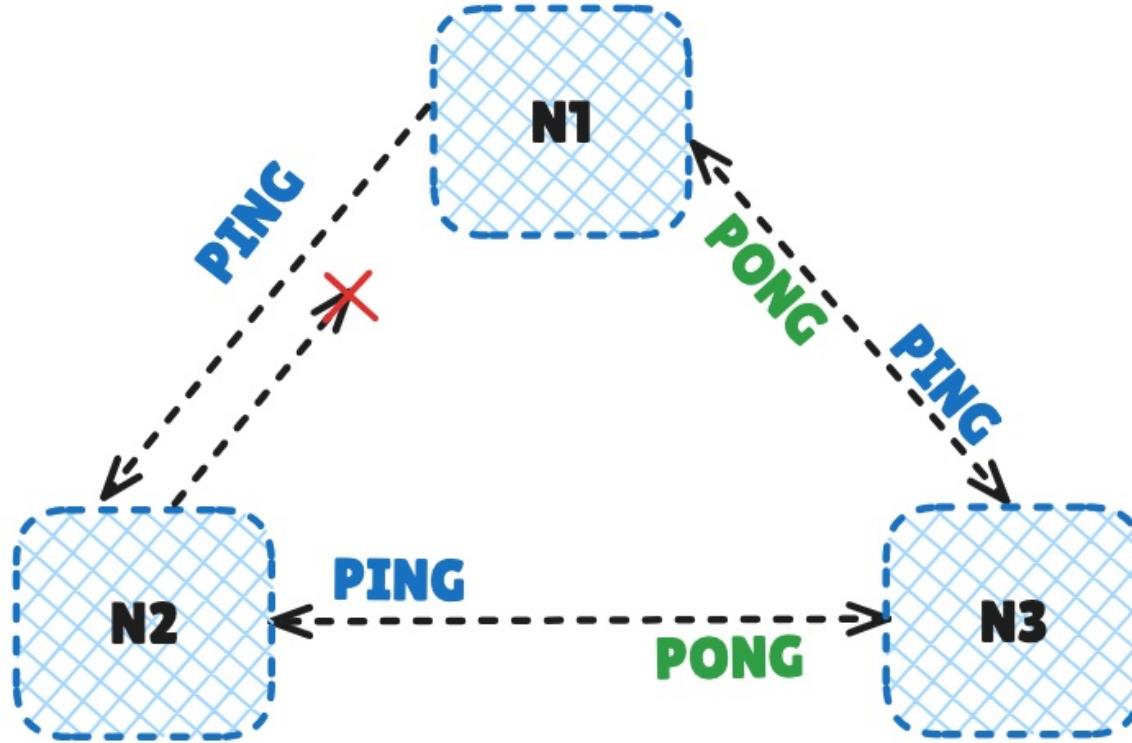
# Failure Detection – PING/PONG

- ping-sent
- pong-received
- Partial failure – node-timeout / 2
- Complete failure (Quorum) – node-timeout
- node-timeout is configurable

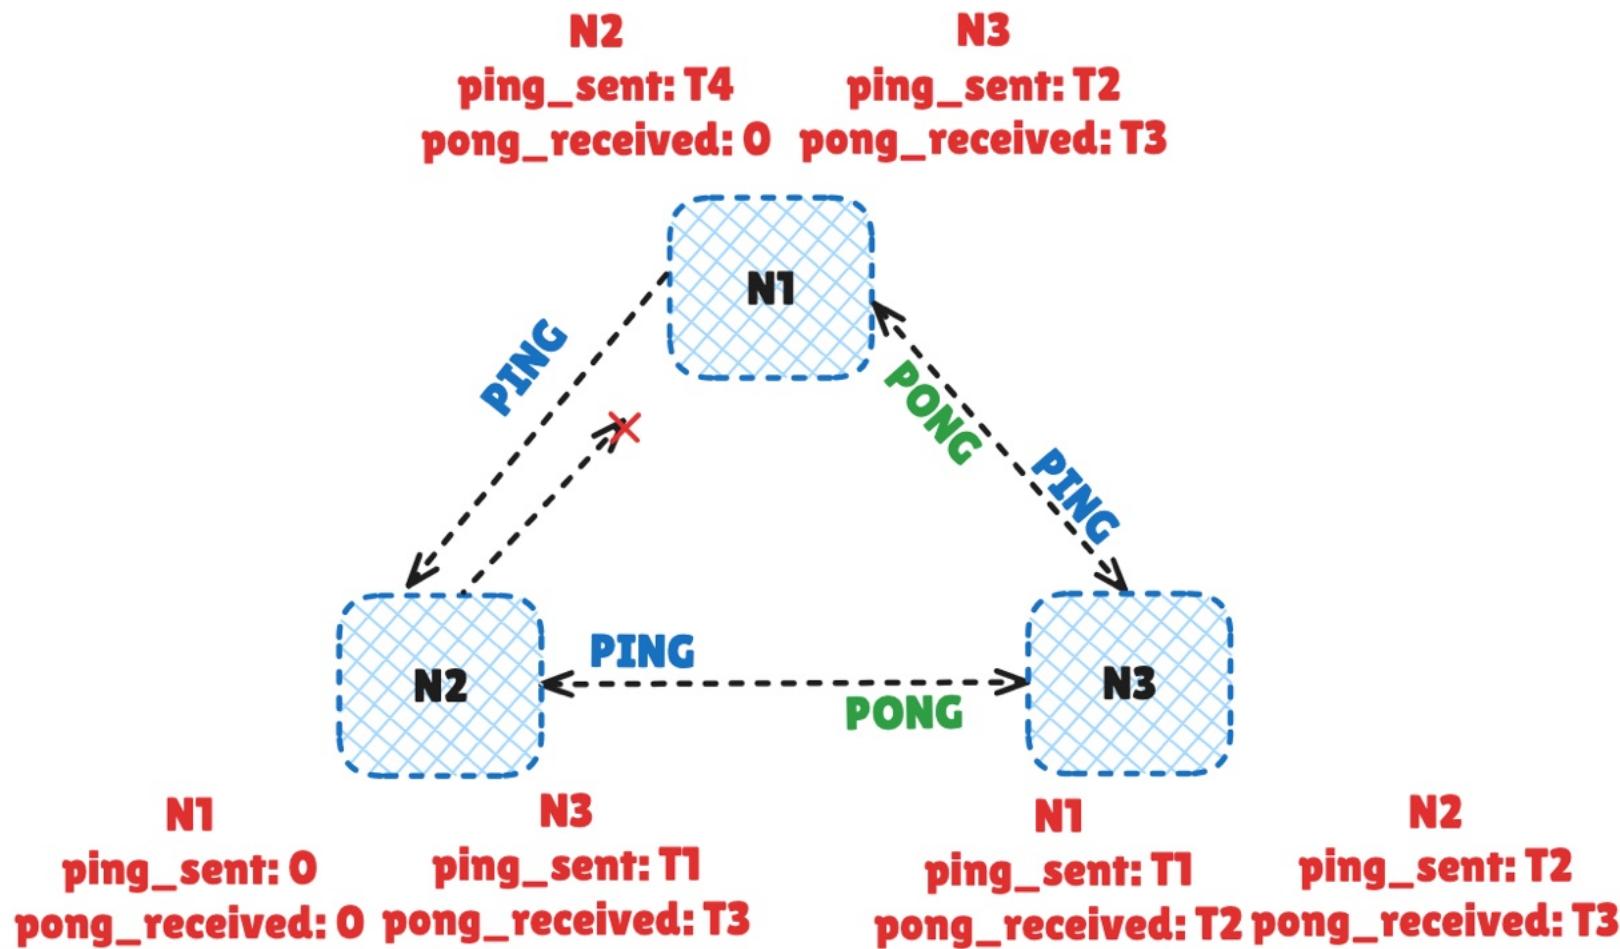
# Failure Detection – Healthy State



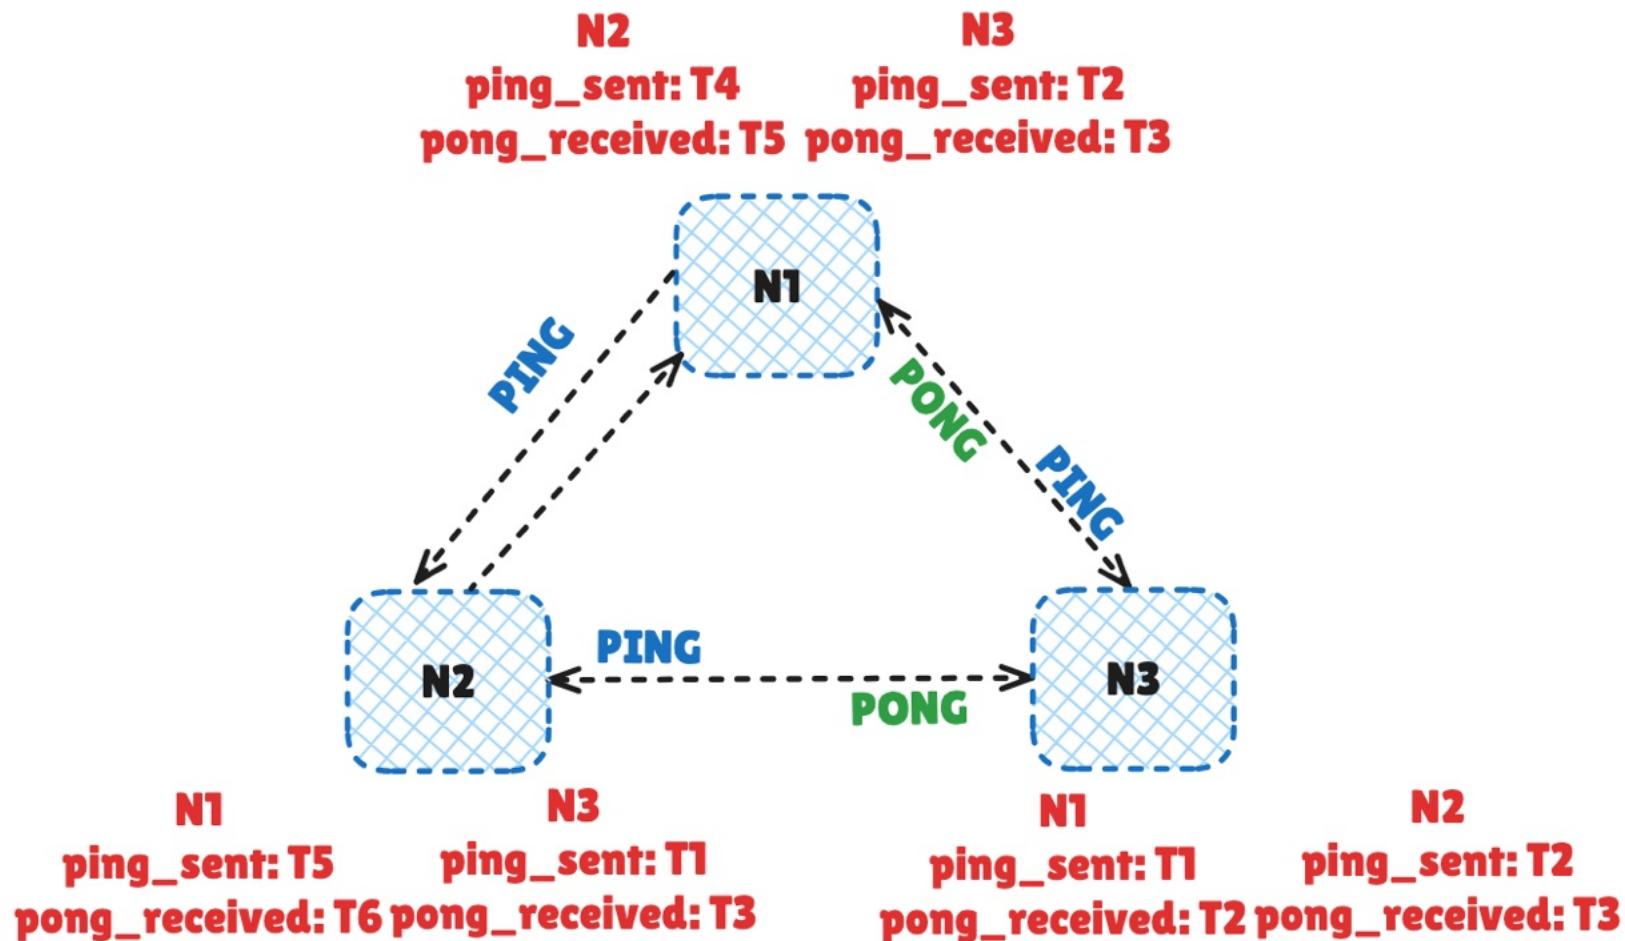
# Failure Detection – Partial failure



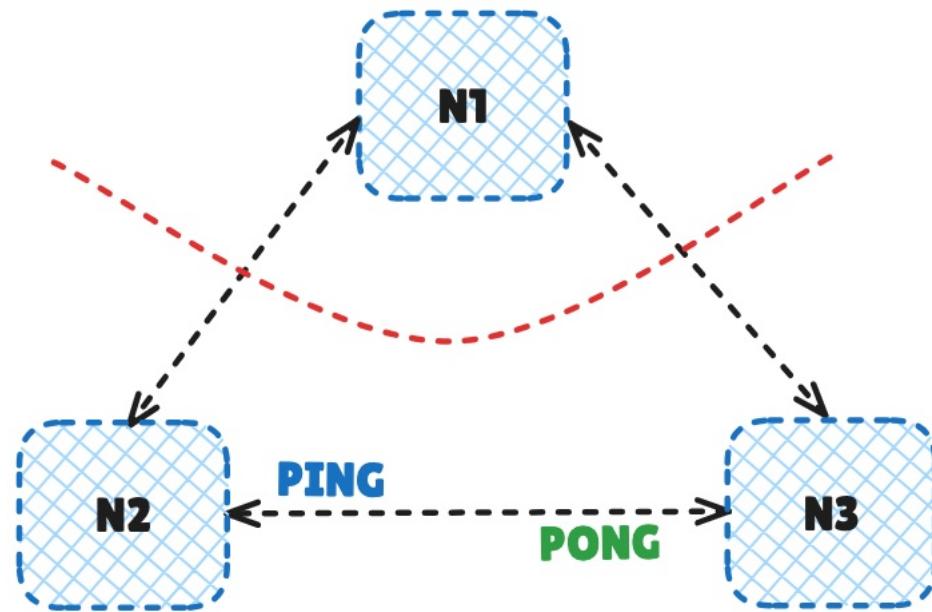
# Failure Detection – Partial failure



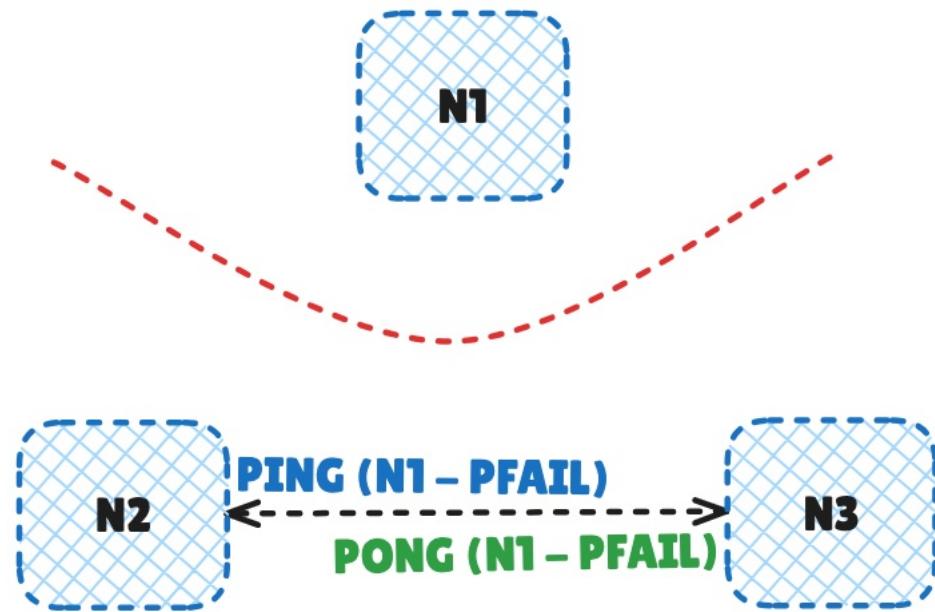
# Failure Detection – Healthy



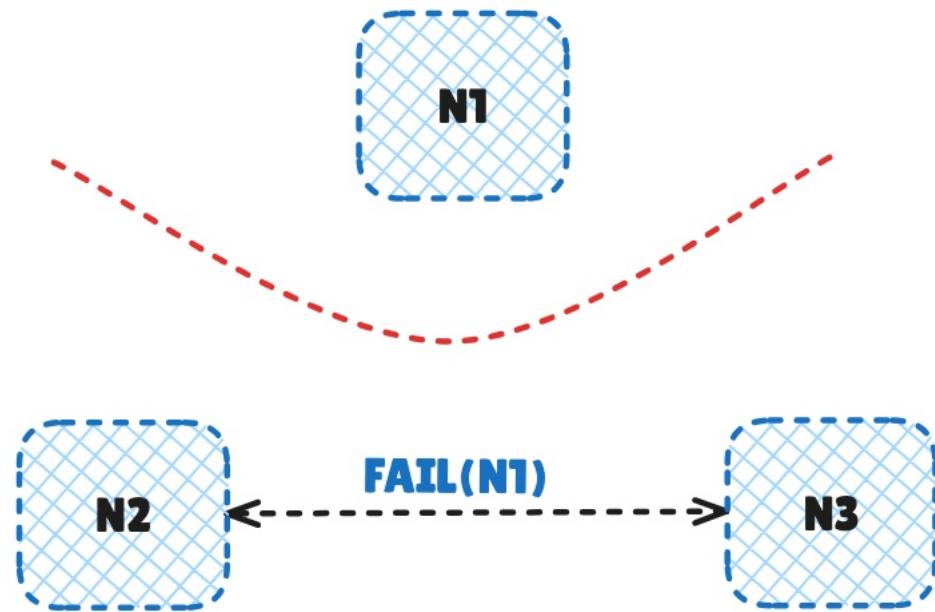
# Failure Detection - Partition



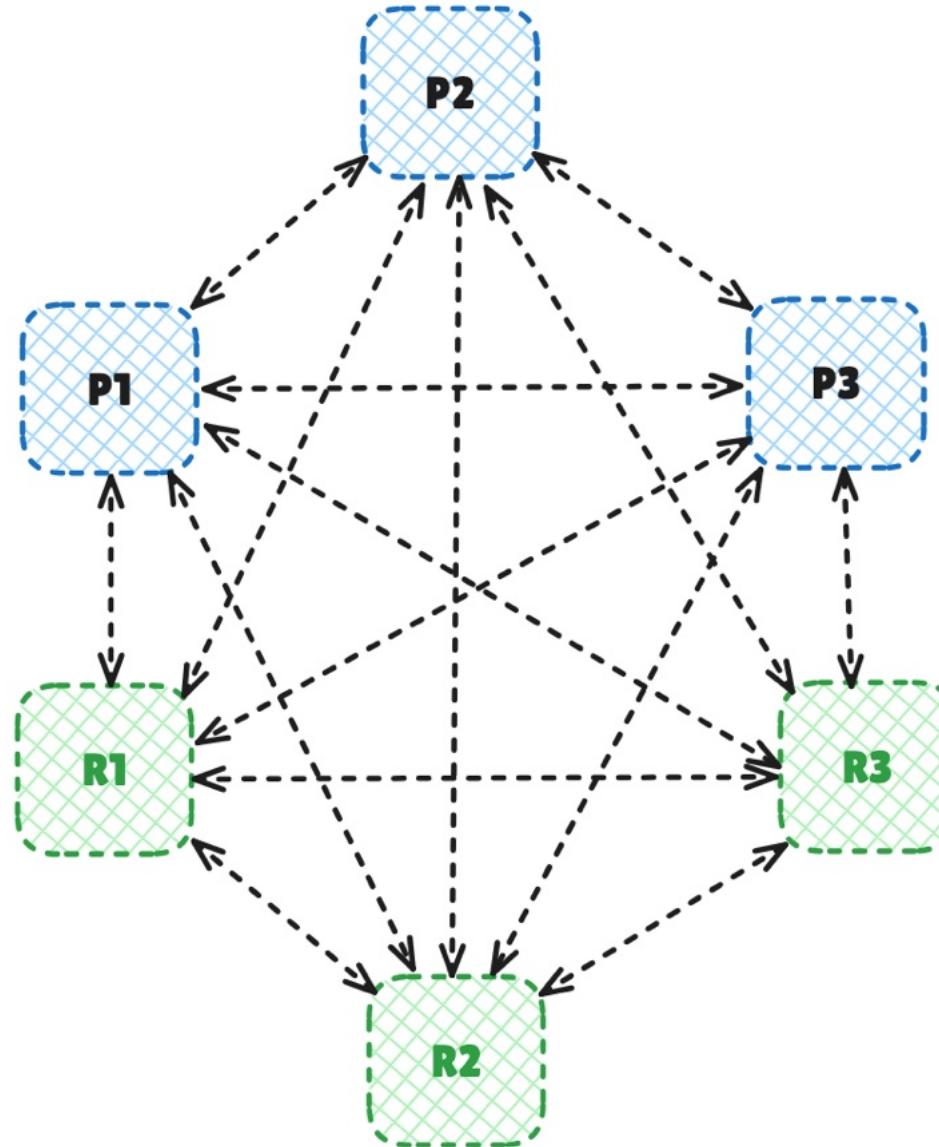
# Failure Detection – Time out gossip



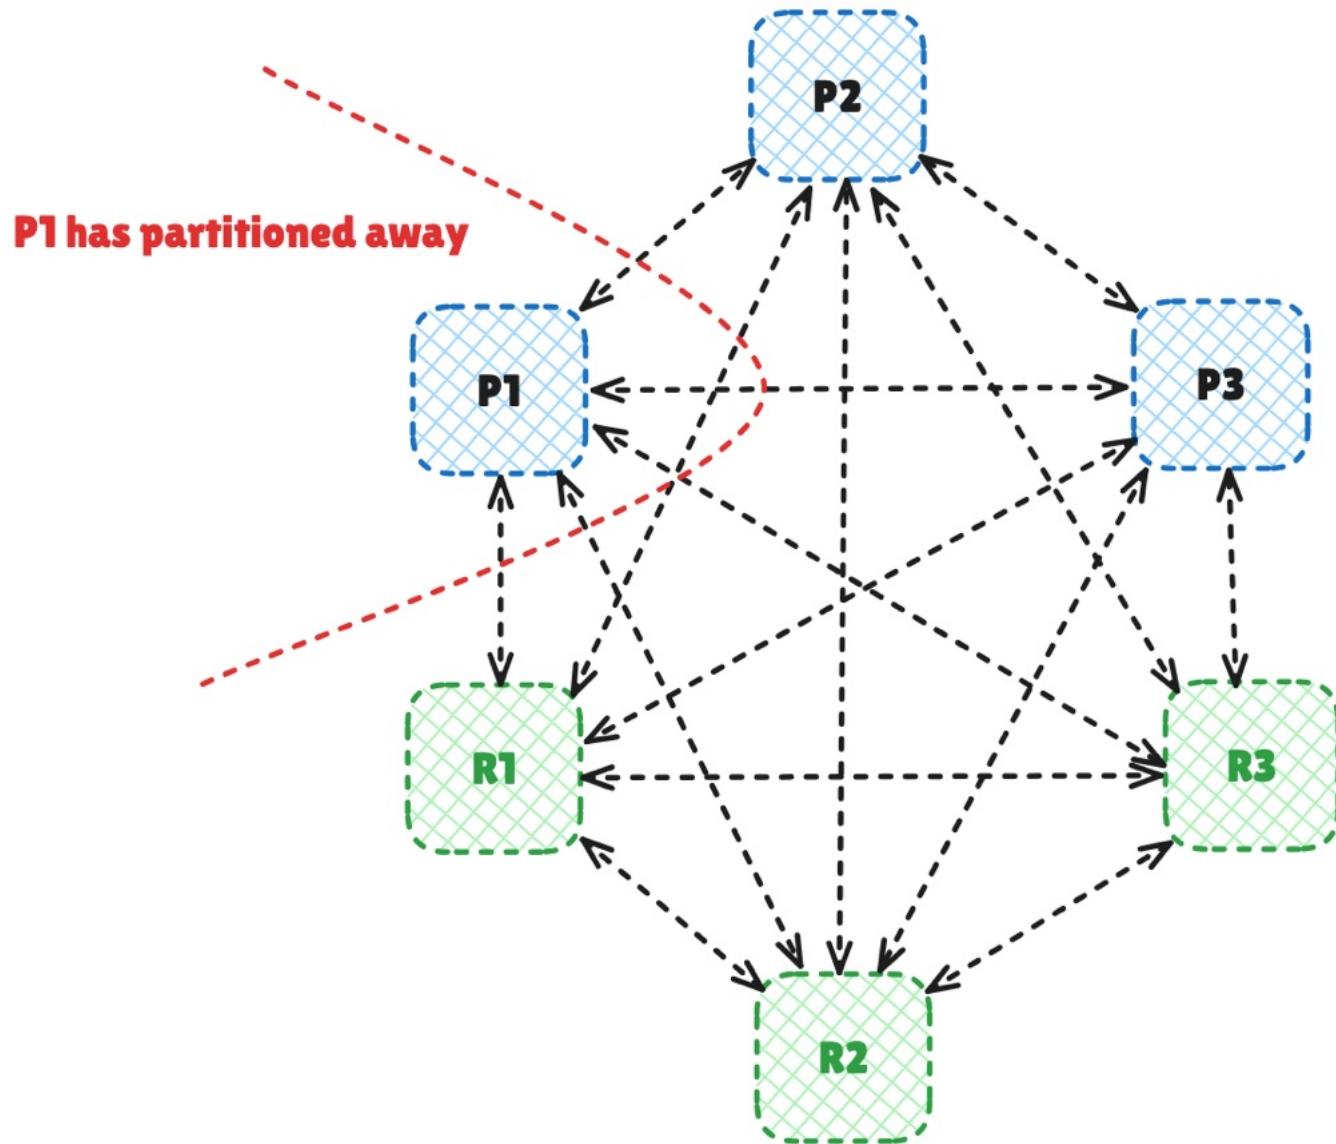
# Failure Detection - Broadcast



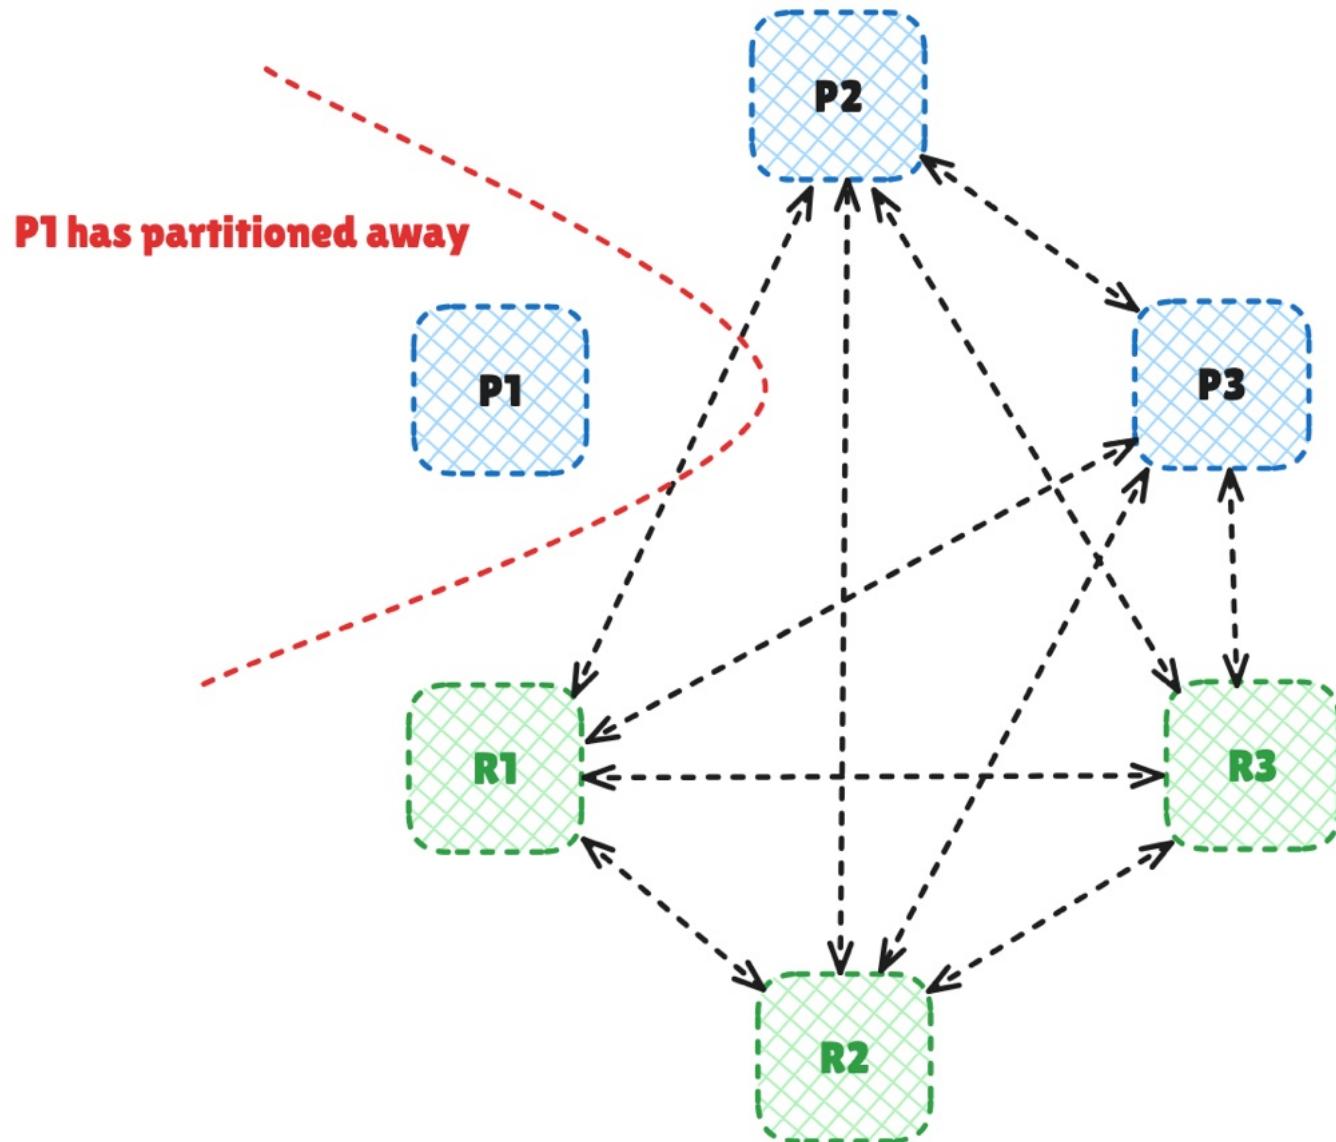
# Failover



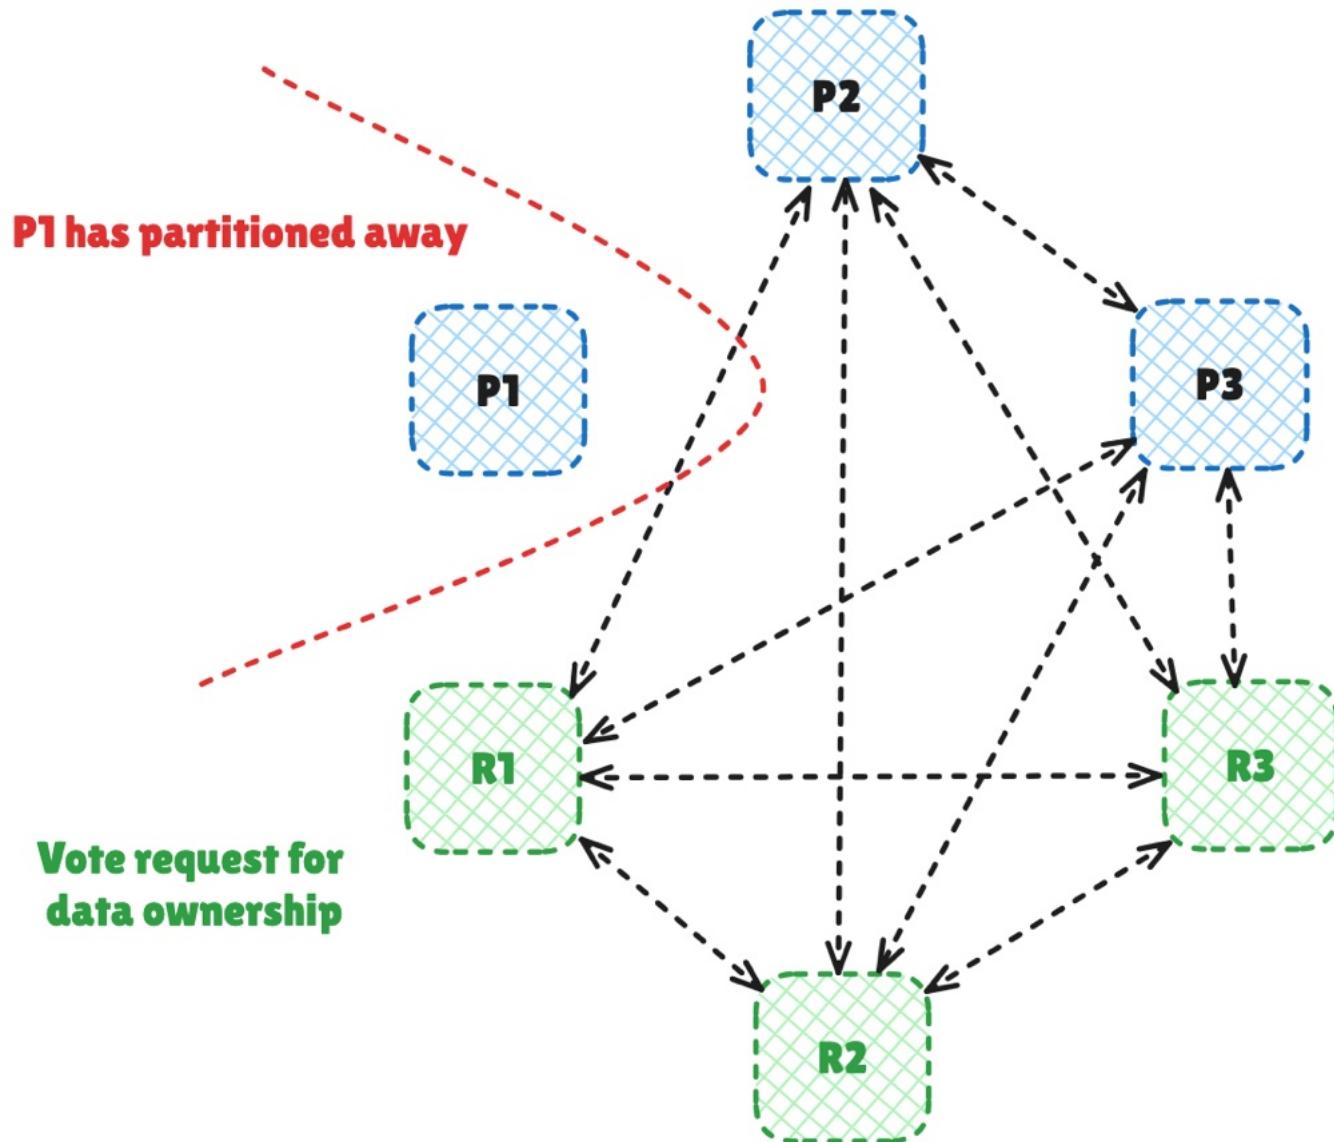
# Failover – Failure detection



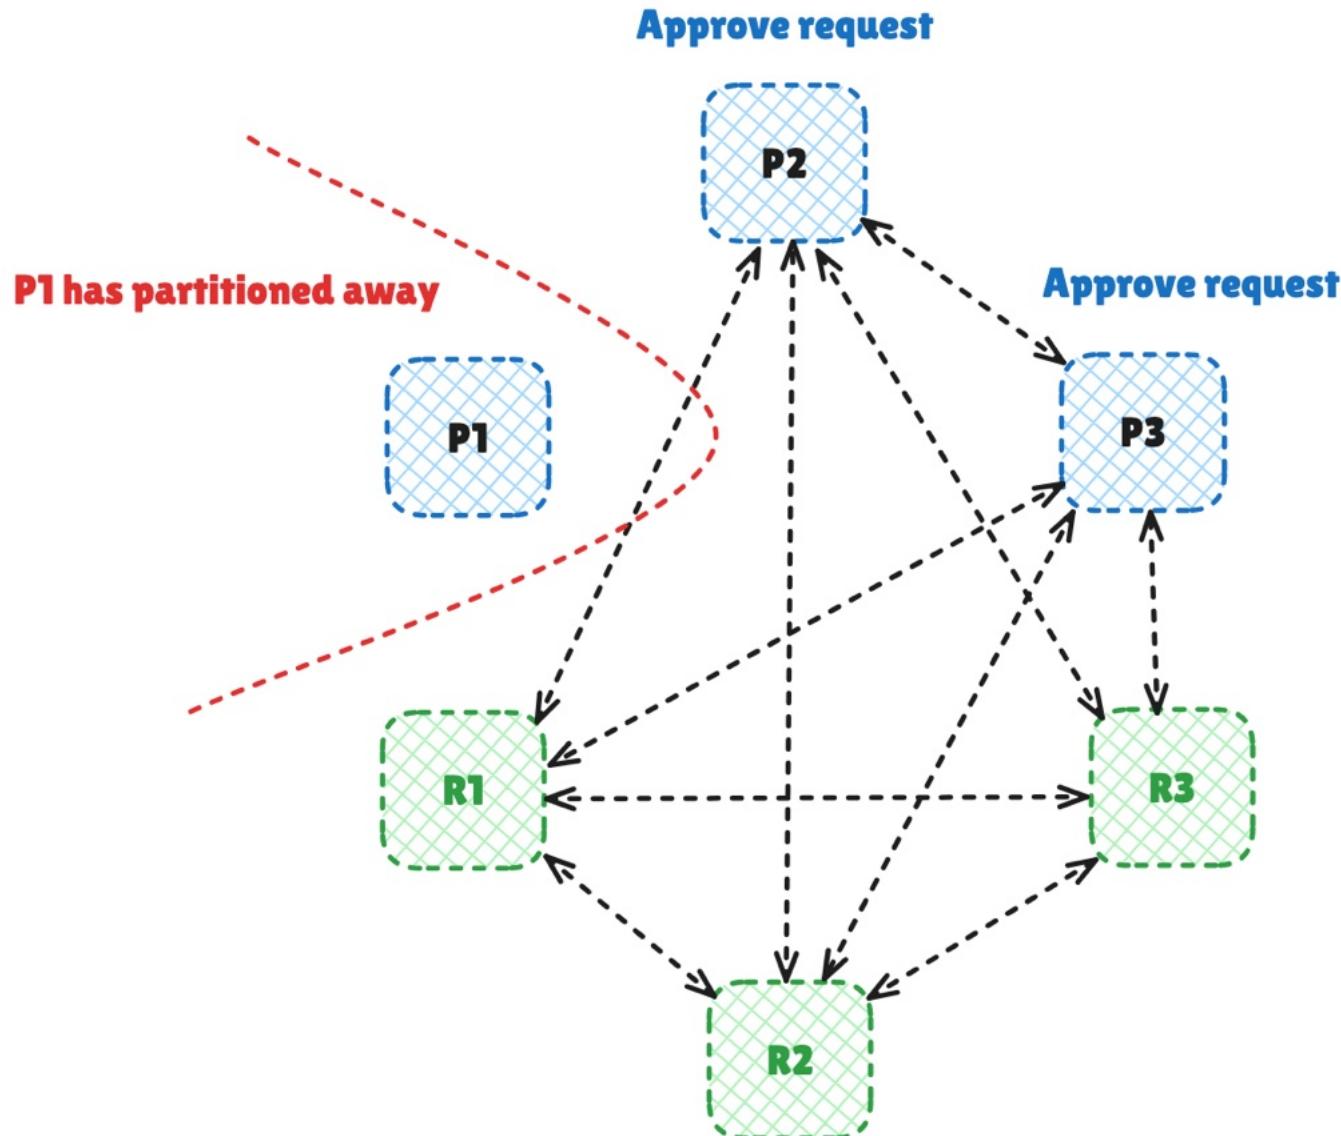
# Failover – Link disconnection



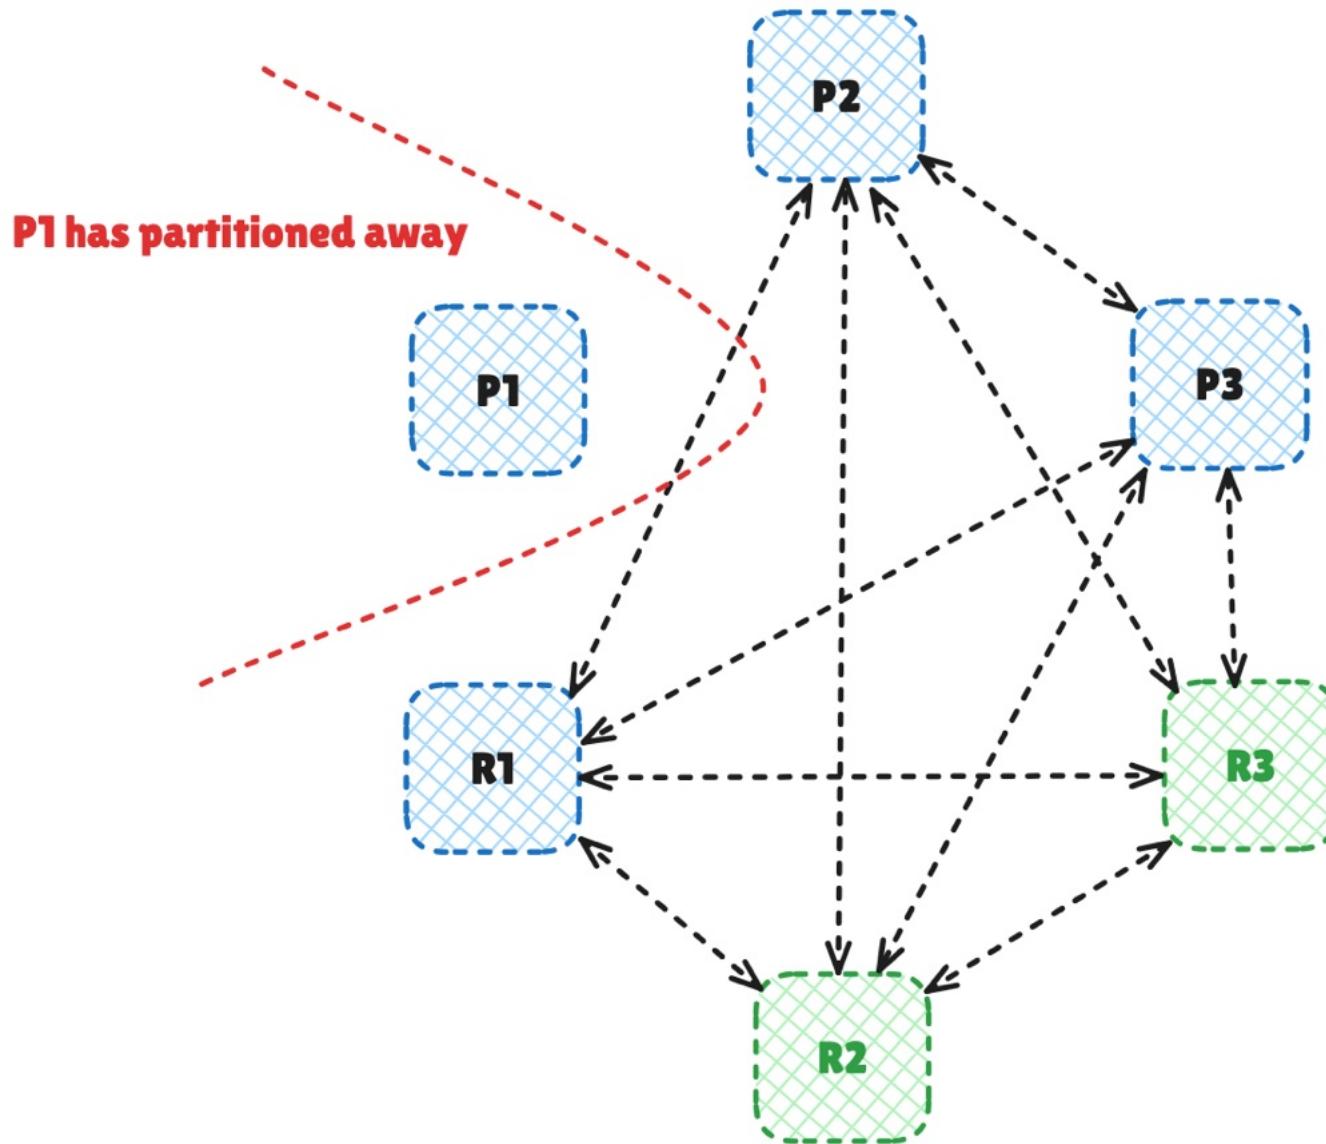
# Failover – Replica promotion request



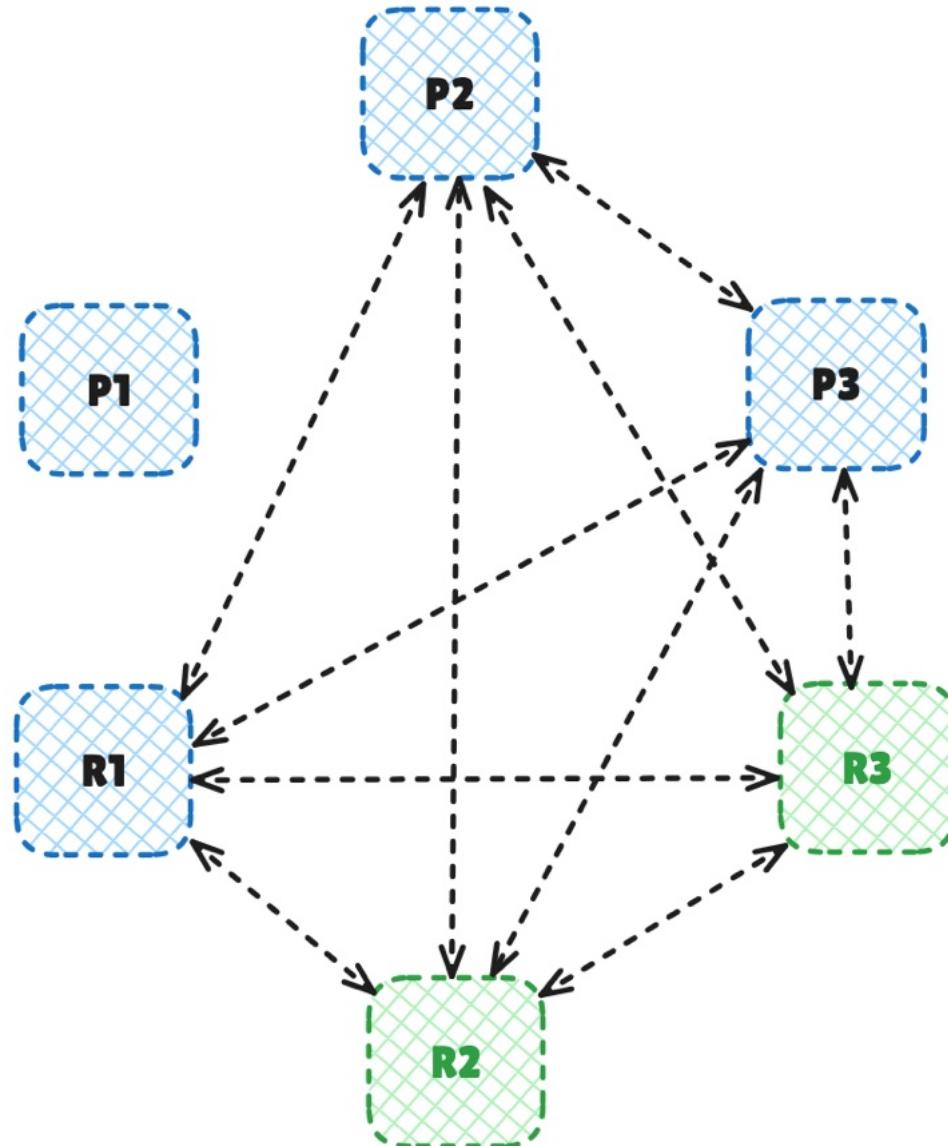
# Failover – Replica promotion request



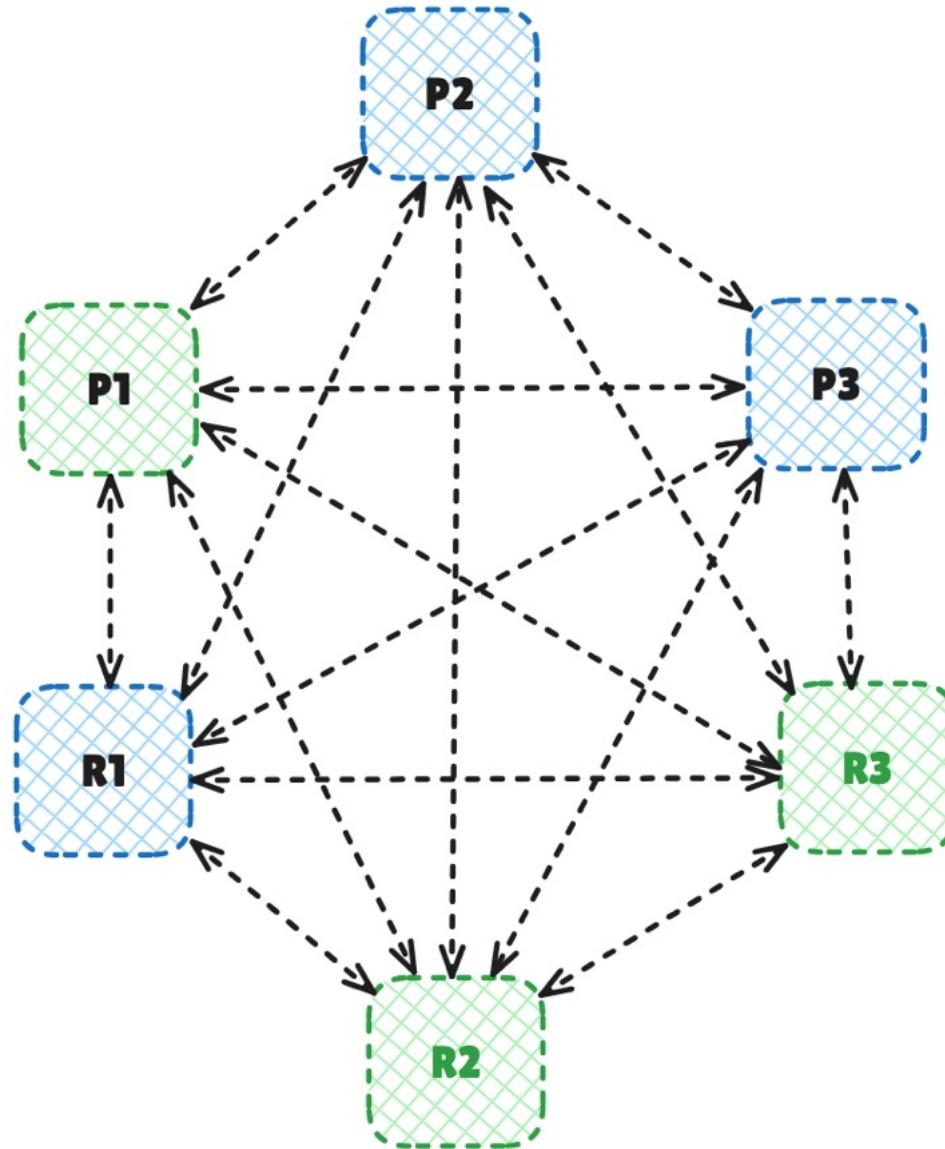
# Failover – Replica promoted



# Failover – Partition heal start



# Failover – New topology



# Conflict Resolution - Epoch

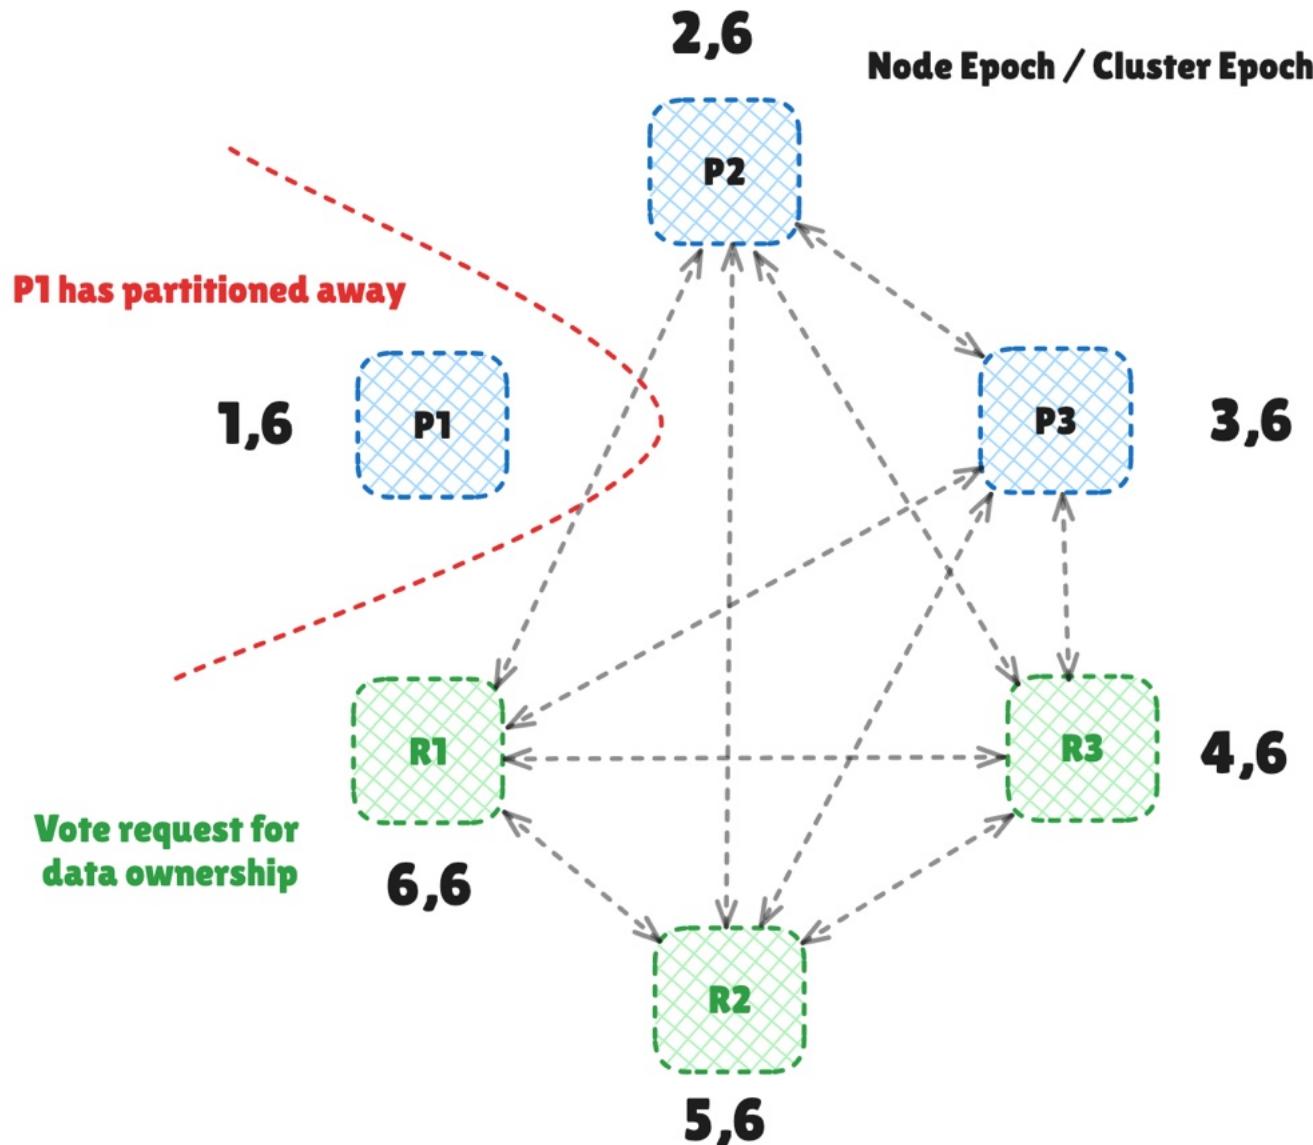
- Decentralized cluster – no single source of truth
- Monotonically increasing counter
- Conflict resolution
  - config epoch, node id to break ties safely



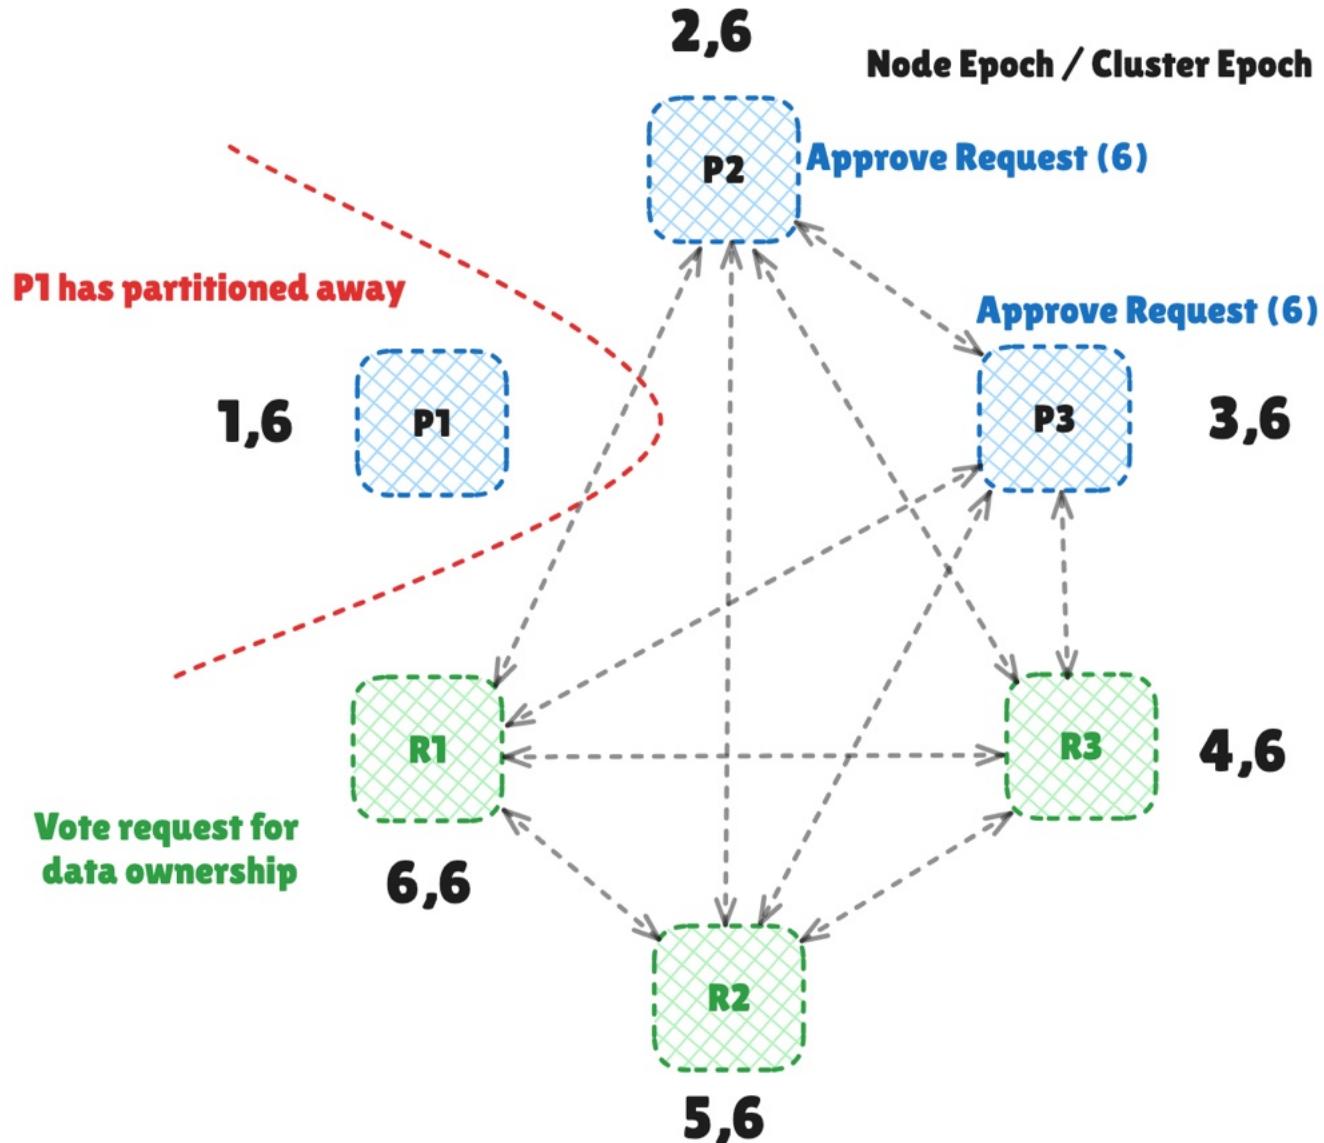
# Conflict Resolution – Epoch Types

- Node level epoch – config epoch
  - Per node epoch indicating **authority over slots**
- Cluster wide epoch – cluster epoch
  - Global for the **entire cluster**
  - Incremented during failovers to assign unique epochs

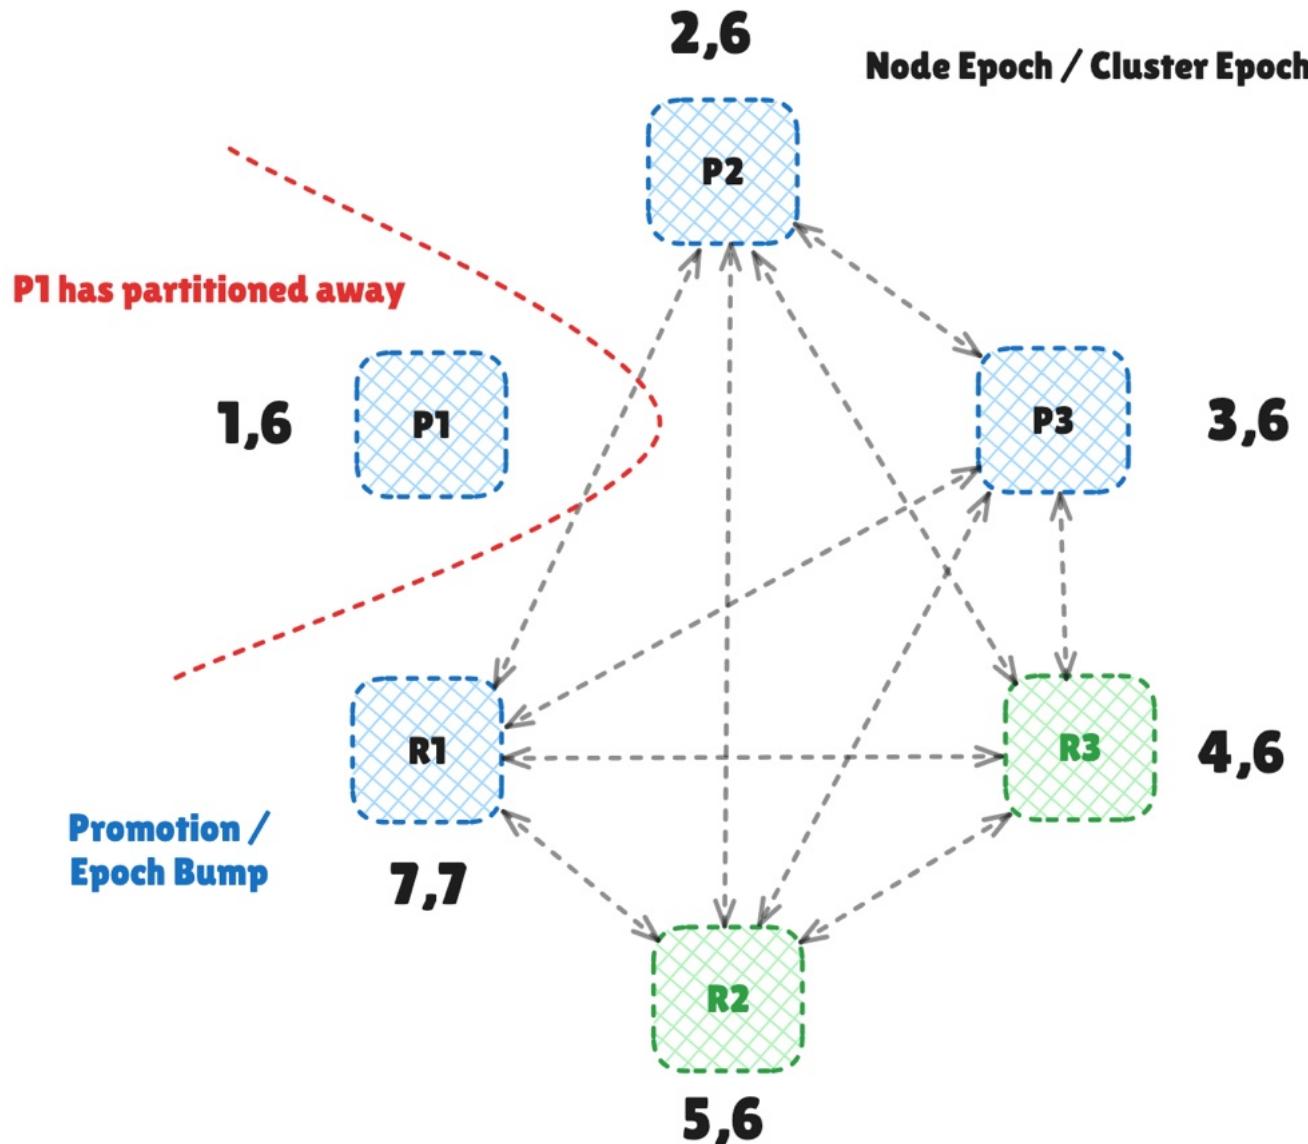
# Conflict Resolution - Epoch



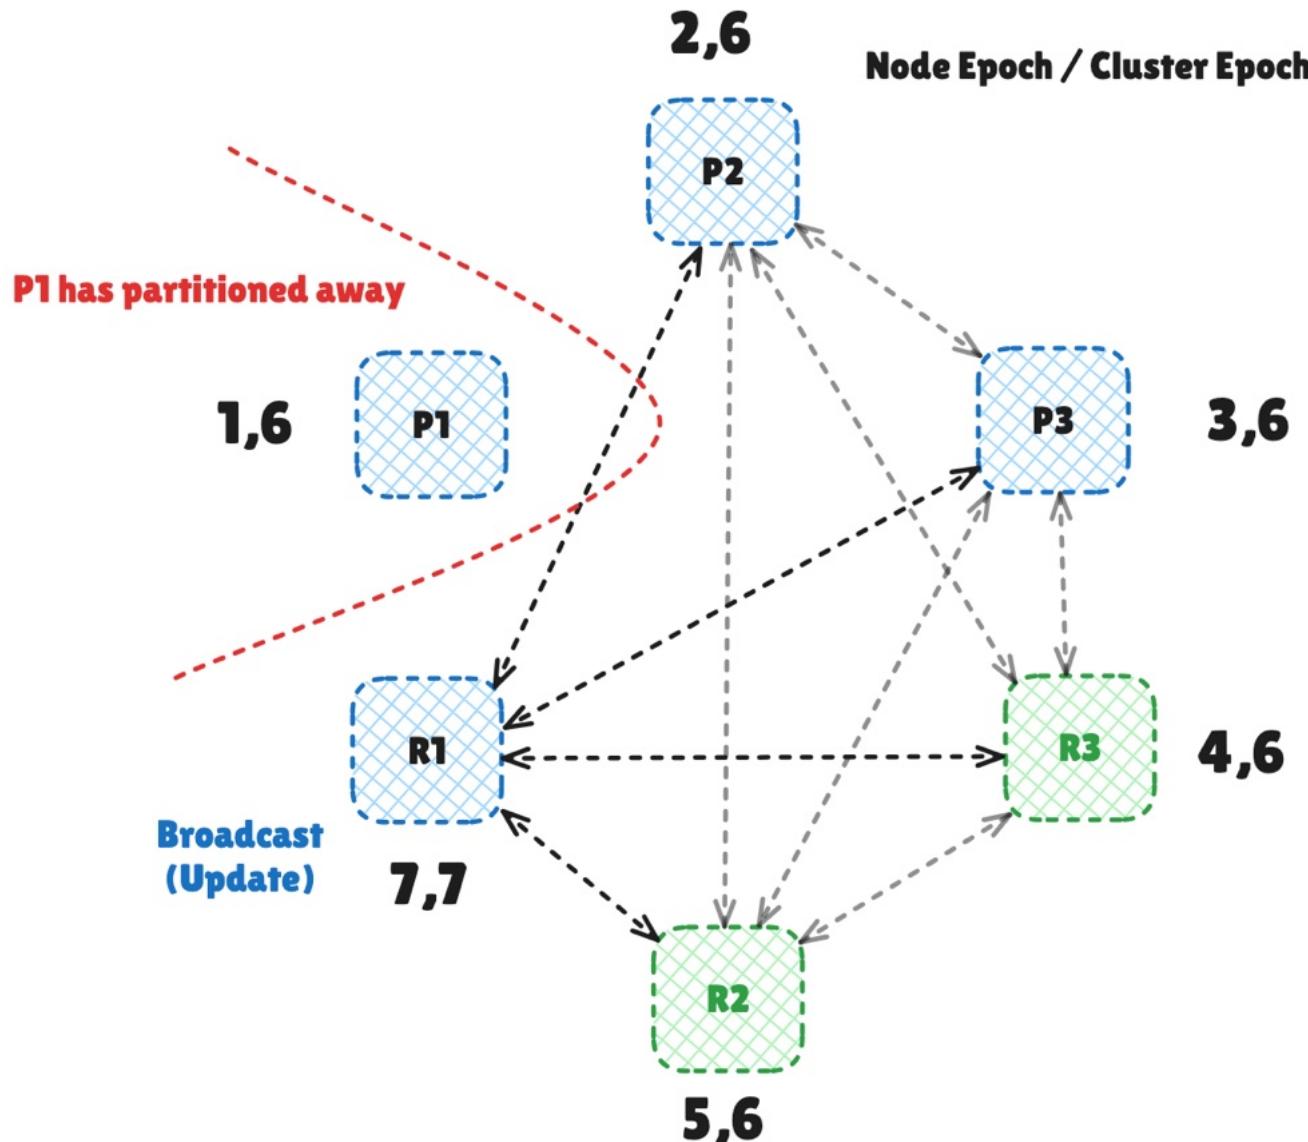
# Epoch – Vote Request



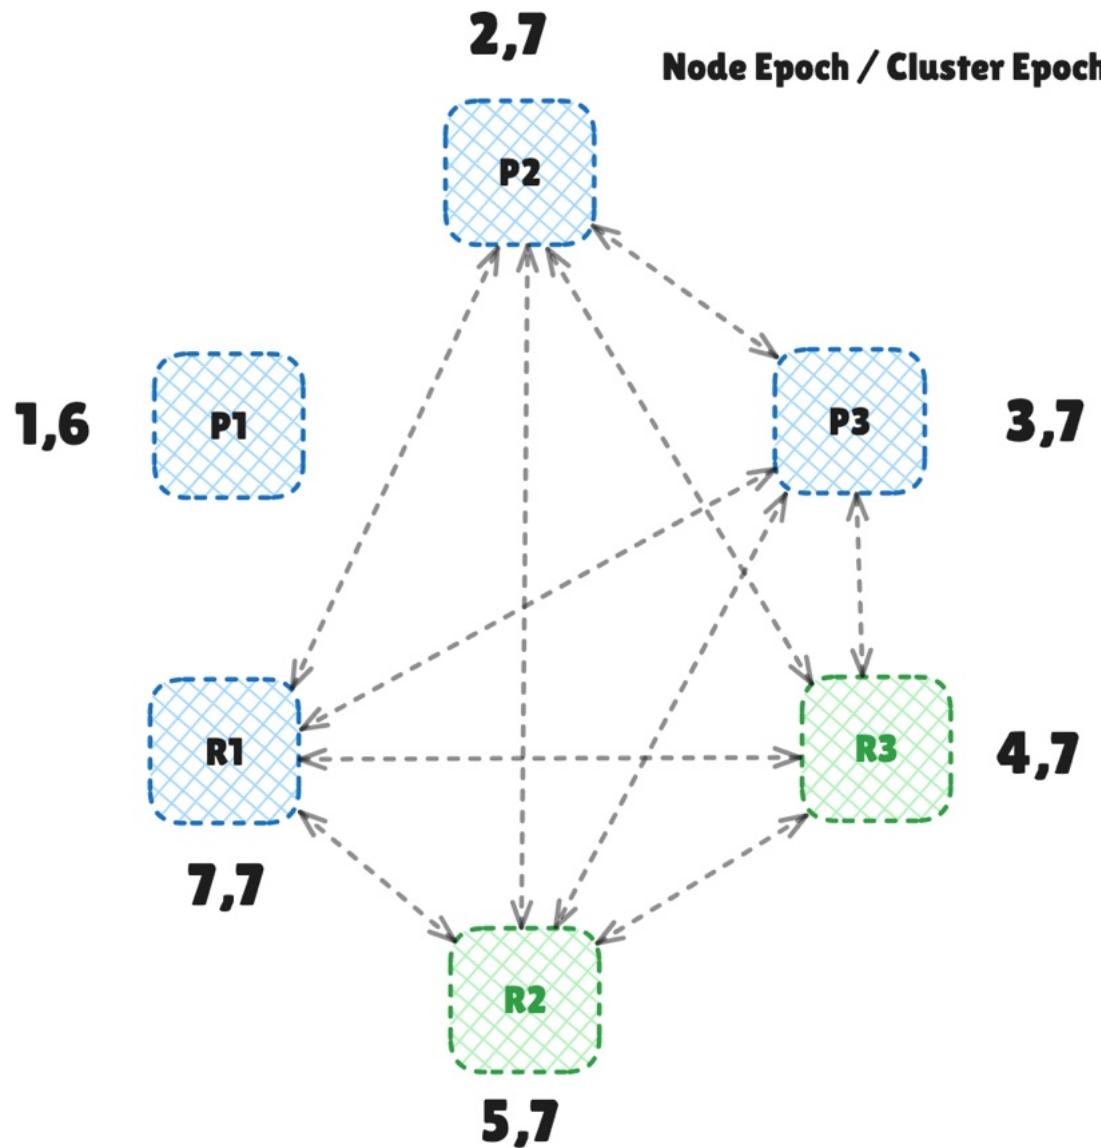
# Epoch Bump



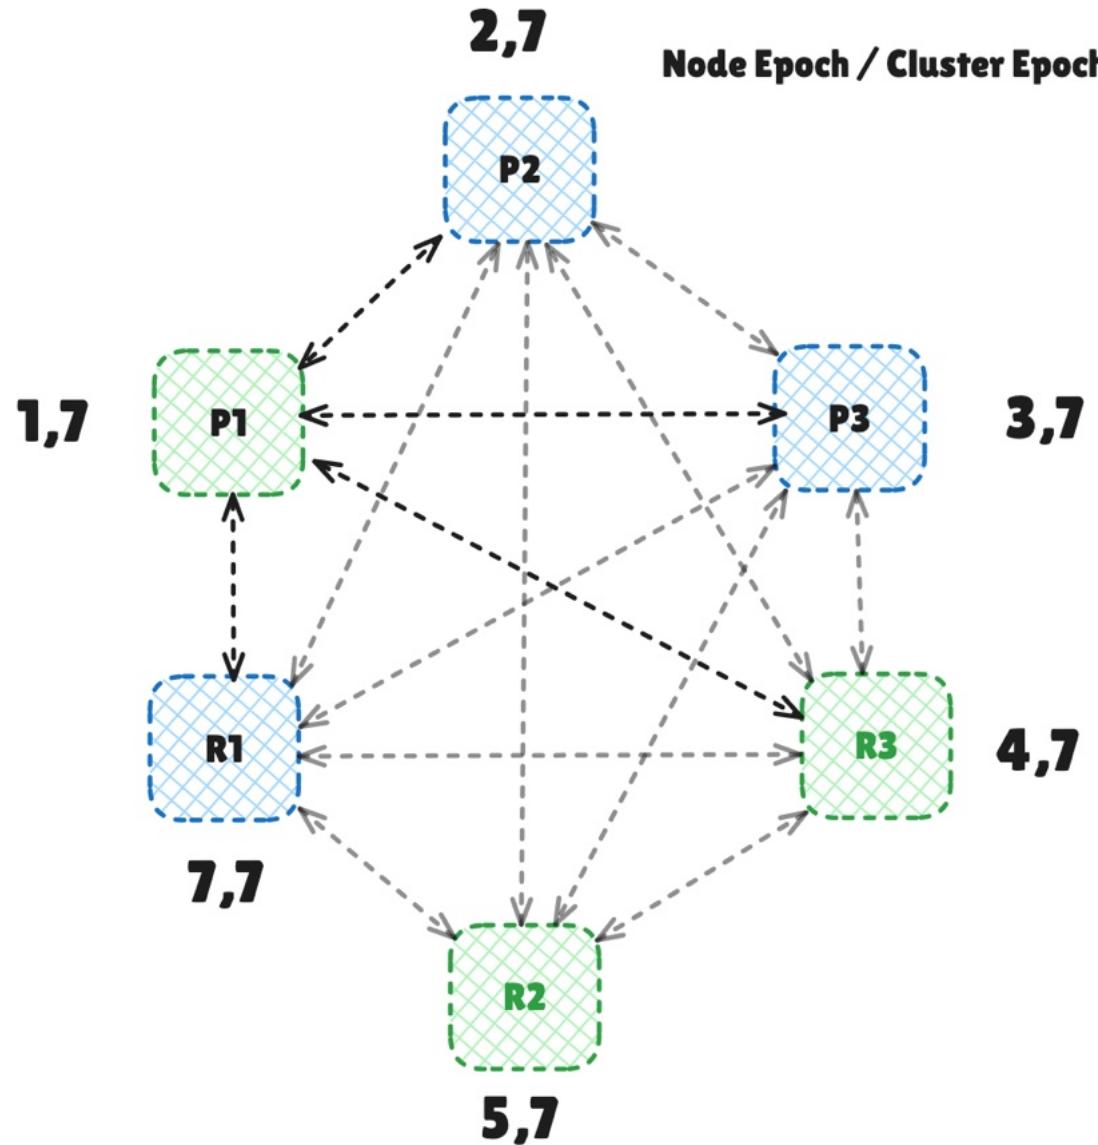
# Epoch Bump - Broadcast



# Epoch Bump – New Topology



# Epoch Bump – Conflict Resolution



# Valkey Clustering improvements

- Fast failover with multiple primary failures
- Light message header type (~30 bytes)
- Connection rate limiting

# Can it scale?

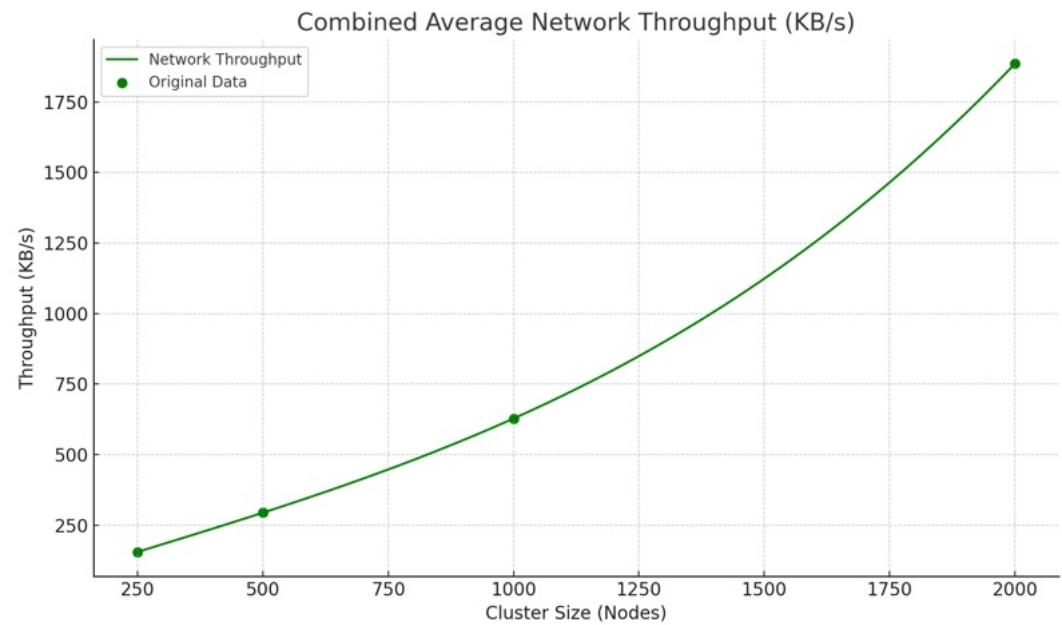
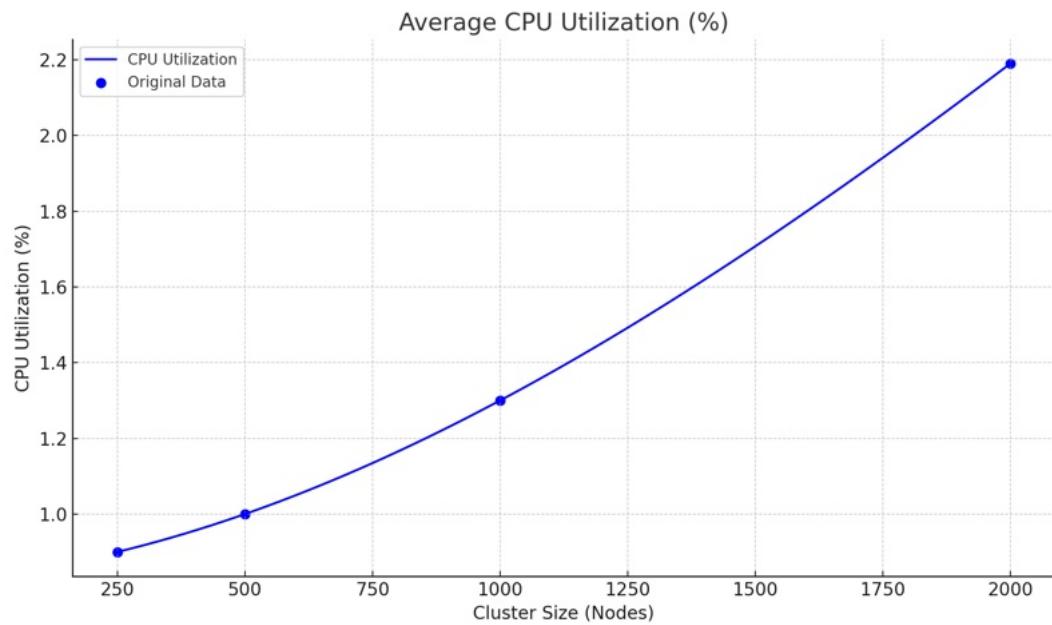
- Peer to peer health detection, does it scale?
- Message transfer rate is high during ideal state ?
- Too many votes (quorum of primaries) for failover ?

# Benchmark - Results

- Scales well upto 2000 nodes cluster
- Node timeout - 15 seconds
- 10% node failure detection and failover time ~ 20 seconds

```
127.0.0.1:6379> CLUSTER INFO
cluster_state:ok
cluster_slots_assigned:16384
cluster_slots_ok:16384
cluster_slots_pfail:0
cluster_slots_fail:0
cluster_nodes_pfail:0
cluster_nodes_fail:0
cluster_voting_nodes_pfail:0
cluster_voting_nodes_fail:0
cluster_known_nodes:1998
cluster_size:666
cluster_current_epoch:1997
cluster_my_epoch:180
cluster_stats_messages_ping_sent:1442617
cluster_stats_messages_pong_sent:1461312
cluster_stats_messages_meet_sent:671
cluster_stats_messages_sent:2904600
cluster_stats_messages_ping_received:1457438
cluster_stats_messages_pong_received:1448050
cluster_stats_messages_meet_received:934
cluster_stats_messages_received:2906422
total_cluster_links_buffer_limit_exceeded:0
127.0.0.1:6379> █
```

# Benchmark - Results



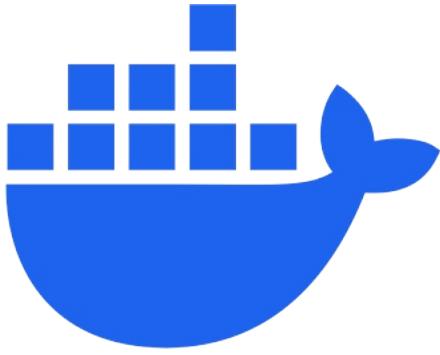
# What's next?

- Reduce steady state CPU utilization
- Offload cluster message processing
- Additional observability information
- Fuzzy testing

# Get involved in the project



[valkey.io](https://valkey.io)



[valkey/valkey](https://github.com/valkey/valkey)



[github.com/valkey-io](https://github.com/valkey-io)

# Thank you!

**Harkrishn Patro**

Software Engineer – AWS

[github.com/hpatro](https://github.com/hpatro)

[linkedin.com/in/harkrishn-patro](https://linkedin.com/in/harkrishn-patro)