

August 28

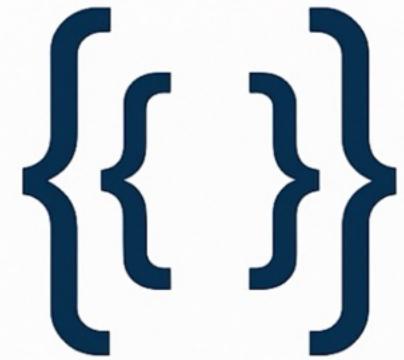
K E Y S P A C E
Amsterdam

valkey-bundle: One stop shop for real-time applications

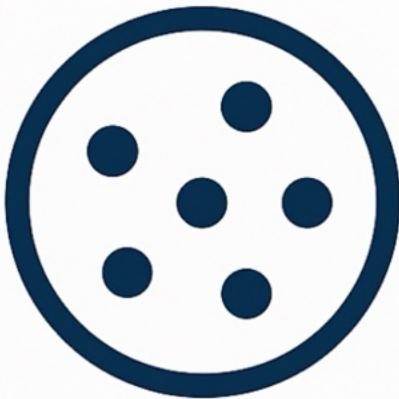
Roberto Luna Rojas
Sr Developer Advocate, Valkey



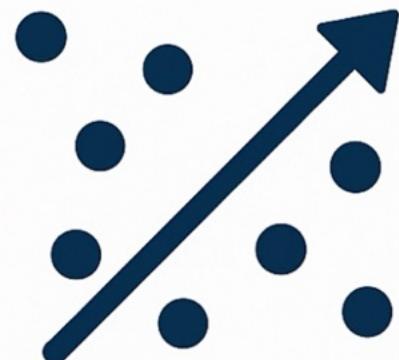
Valkey-Bundle: Building Modern Low-Latency Applications



JSON



Bloom Filter

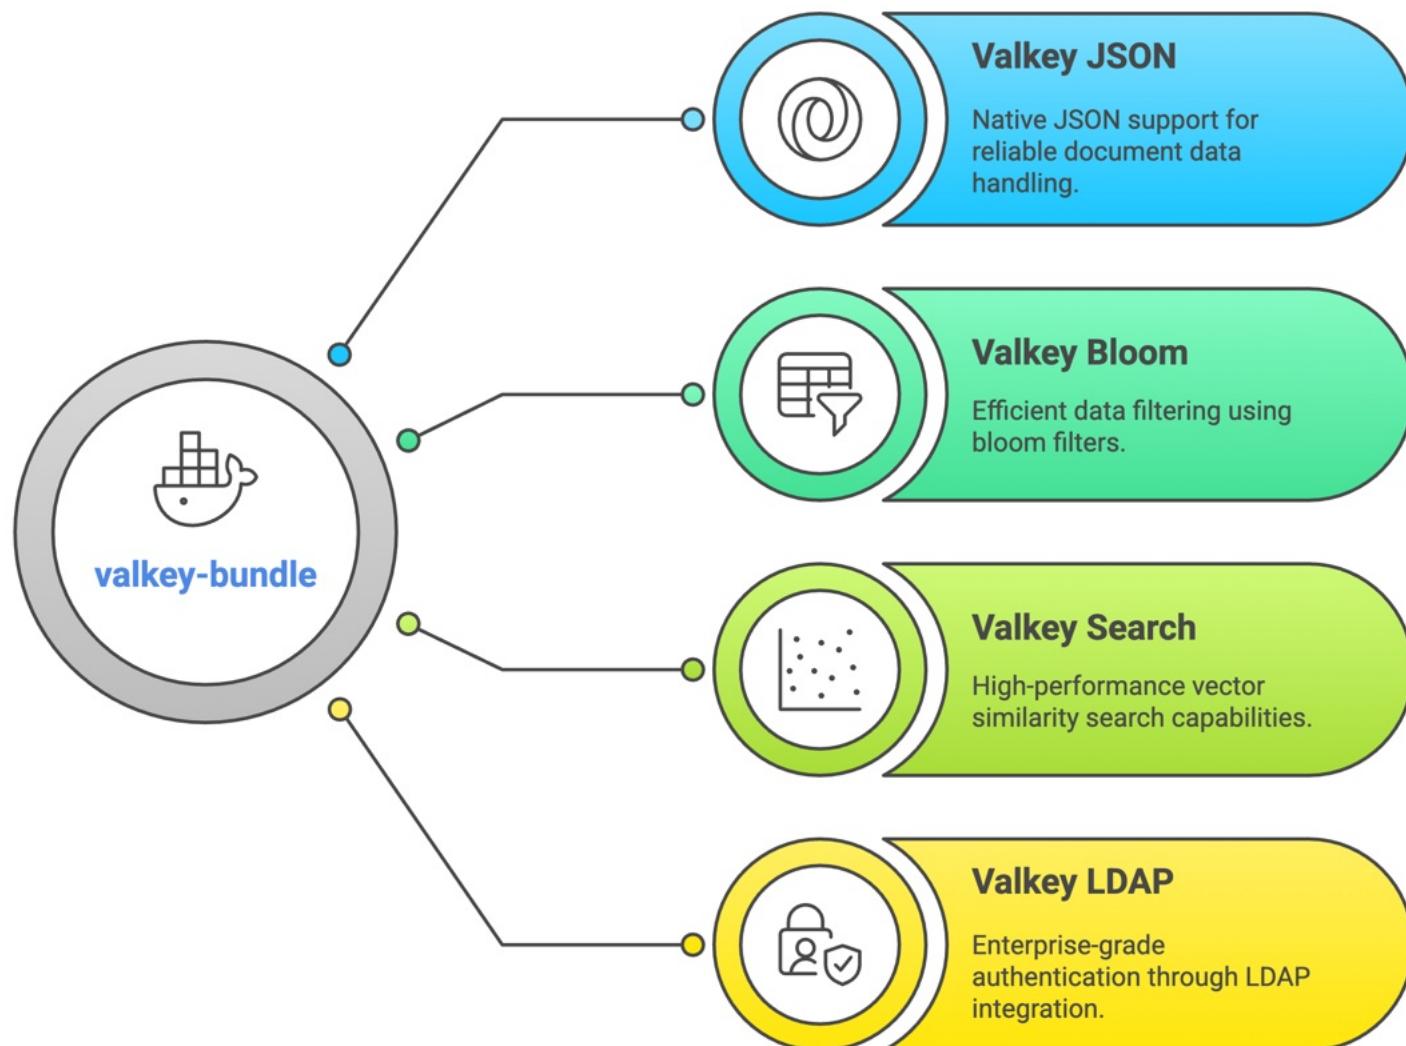


Vector Search



LDAP Auth

valkey-bundle: One stop shop for real-time applications

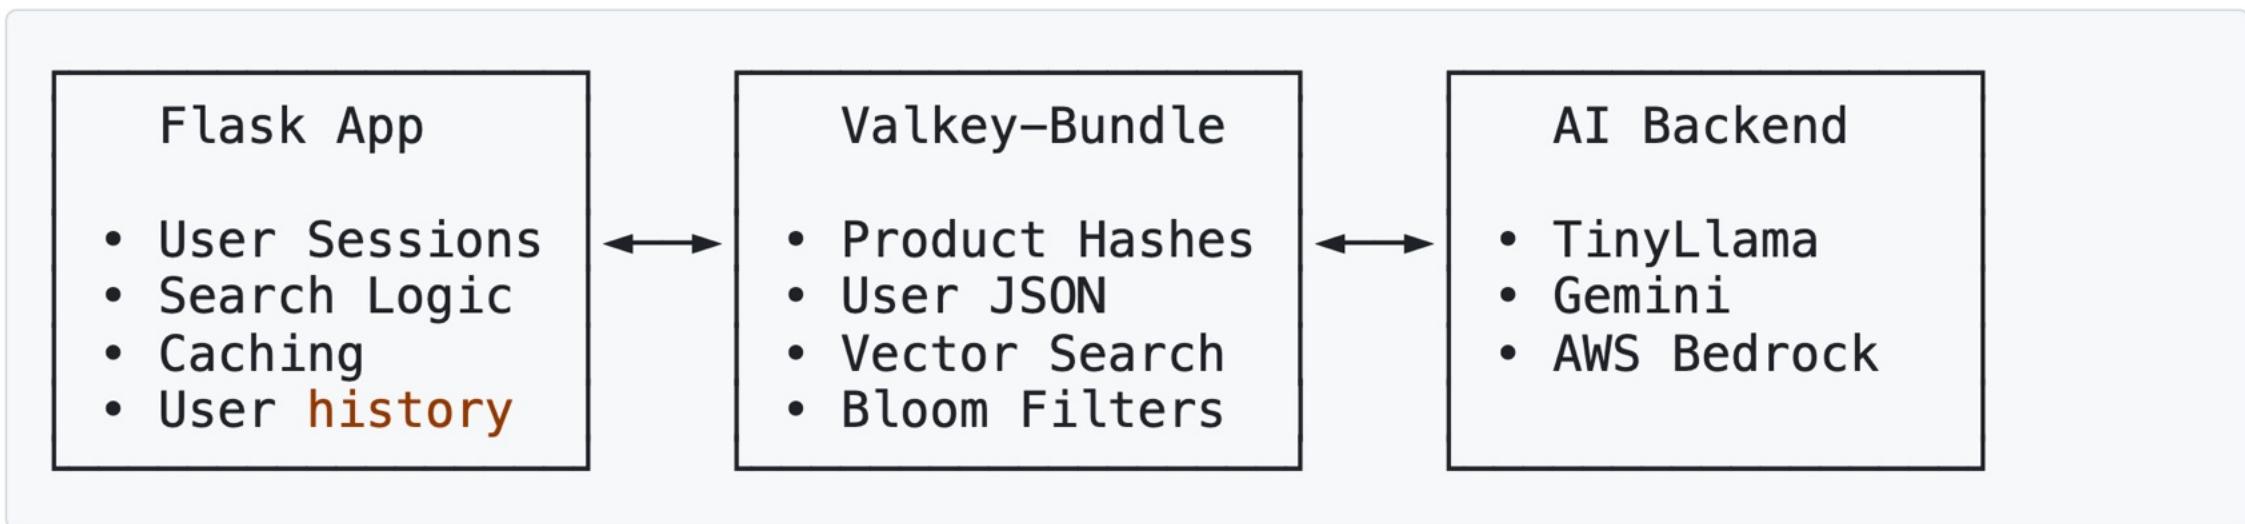


Application Architecture Overview

Our demonstration application is a personalized product search system that combines:

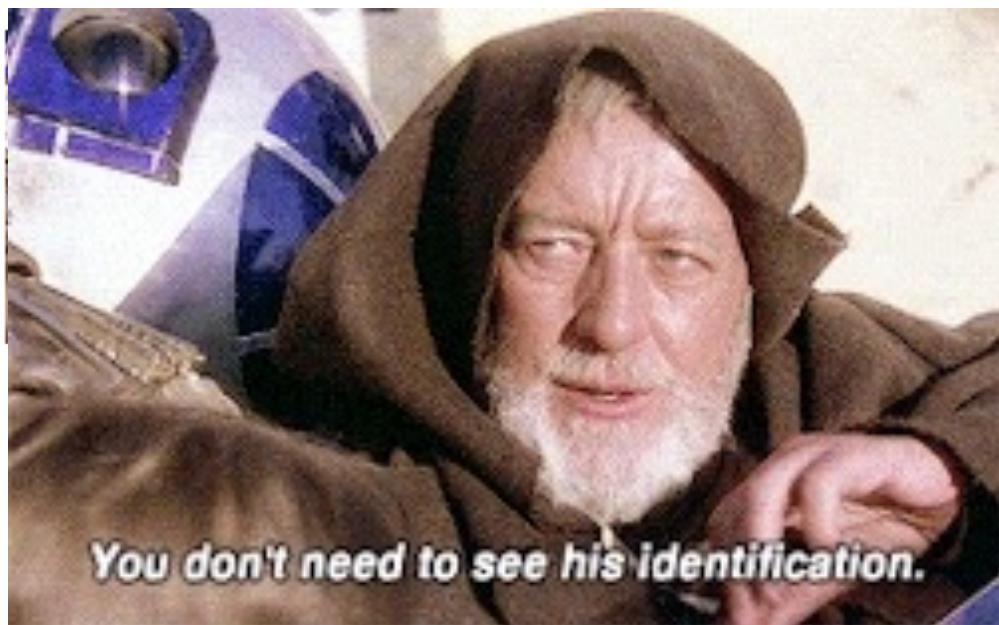
- **AI-Powered Personalization:** Using local TinyLlama, Google Gemini, or AWS Bedrock Nova Pro
- **Hybrid Search:** Traditional keyword filtering + vector similarity search
- **Real-time Recommendations:** Personalized product descriptions generated on-demand
- **Efficient Tracking:** User behavior monitoring with probabilistic data structures
- **High-Performance Caching:** LLM response caching to minimize compute costs

Application Architecture Overview (cont)



LDAP Integration

Session: Deploying Valkey at Enterprise level with LDAP authentication and auditing at
13:45



You don't need to see his identification.

Component 1: Product Storage with Valkey Hashes

The Challenge

Storing complex product data with multiple attributes while maintaining fast access patterns for search and retrieval.

The Solution: Valkey Hashes

Products are stored as hash structures, providing efficient field-level access.

Valkey Hashes Python example

```
# Product storage structure
product_data = {
    'id': 12345,
    'name': 'Wireless Bluetooth Headphones',
    'brand': 'TechAudio',
    'main_category': 'Electronics',
    'sub_category': 'Audio',
    'price': 89.99,
    'rating': 4.5,
    'review_count': 1247,
    'search_tags': 'wireless,bluetooth,headphones,audio,music',
    'region': 'NA',
    'embedding': <384/768/1024-dimensional vector bytes>
}

# Stored as: HSET product:12345 field1 value1 field2 value2 ...
```

Why Hashes?

- **Memory Efficient:** Optimized storage for objects with multiple fields
- **Atomic Operations:** Update individual fields without affecting others
- **Fast Access:** $O(1)$ field retrieval and updates
- **Structured Data:** Natural mapping to application objects

Embedding Integration

Each product includes a vector embedding generated by:

- **Local Mode:** sentence-transformers (384 dimensions)
- **Google Cloud:** Vertex AI text-embedding-004 (768 dimensions)
- **AWS Bedrock:** Titan Text Embeddings v2 (1024 dimensions)

The embedding captures semantic meaning of the product for similarity search.

Demo: Valkey Hash for Products

Storing the value for product 123456789

```
HSET product:123456789 'id' 123456789 'name' 'Wireless Bluetooth Headphones' 'brand' 'TechAudio' 'main_category' 'Electronics' 'sub_category' 'Audio' 'price' 89.99 'rating' 4.5 'review_count' 1247 'search_tags' 'wireless,bluetooth,headphones,audio,music' 'region' 'NA'
```

Response:

```
(integer) 10
```

Retrieving the value for product 123456789

```
HGETALL product:123456789
```

Response:

```
1# "id" => "123456789"
2# "name" => "Wireless Bluetooth Headphones"
3# "brand" => "TechAudio"
4# "main_category" => "Electronics"
5# "sub_category" => "Audio"
6# "price" => "89.99"
7# "rating" => "4.5"
8# "review_count" => "1247"
9# "search_tags" => "wireless,bluetooth,headphones,audio,music"
10# "region" => "NA"
```

Or retrieving specific field to reduce latency and network transfer

```
HGET product:123456789 name
```

Response:

```
"Wireless Bluetooth Headphones"
```

Or retrieving multiple cherry picked fields

```
HMGET product:123456789 name price rating
```

Response:

- 1) "Wireless Bluetooth Headphones"
- 2) "89.99"
- 3) "4.5"

Component 2: User Profiles with Valkey-JSON

The Challenge

Storing complex user profiles with nested data structures, purchase history, and dynamic attributes that may evolve over time.

The Solution: Valkey-JSON

Users are stored as native **JSON** documents, enabling rich data structures:

Valkey-JSON Python example

```
# User profile structure
user_profile = {
    "id": "6379",
    "name": "Roberto Luna-Rojas",
    "country": "Mexico ",
    "bio": "Tech enthusiast and early adopter who loves cutting-edge gadgets...",
    "avatar": "data:image/svg+xml;base64,PHN2ZyB2aWV3Qm94PSIwIDAgODAgODAi...",
    "purchase_history": [
        {"product_id": 456, "date": "2024-12-15", "rating": 5, "price": 123.45},
        {"product_id": 789, "date": "2024-11-20", "rating": 4, "price": 234.56}
    ],
    "preferences": {
        "categories": ["electronics", "gaming"],
        "price_range": {"min": 50, "max": 500},
        "brands": ["Apple", "Samsung", "Sony"]
    },
    "embedding": [0.1234, -0.5678, 0.9012, ...] // User preference vector
}

# Stored as: JSON.SET user:6379 $ '{"id":"101","name":"Roberto Luna-Rojas",...}'
```

Why JSON?

- **Flexible Schema:** Easy to add new fields without migration
- **Nested Structures:** Natural representation of complex data
- **Atomic Updates:** Modify specific paths within the document
- **Query Capabilities:** JSONPath queries for complex data retrieval
- **Type Preservation:** Maintains data types (numbers, booleans, arrays)

Demo: Valkey JSON

Store the user JSON document

```
JSON.SET user:6379 $ {"id": 6379, "name": "Roberto Luna-Rivas", "country": "Mexico 🇲🇽", "bio": "Tech enthusiast and early adopter who loves cutting-edge gadgets...", "avatar": "data:image/png;base64,iVBORw0KGgoAAAANSUhEUgAAQAAAABQCAQAAACBQAAAABGdBTUEAALGPC/xhBQAAQ3Nzc3Nz..."}, "purchase history": [{"product id": 456, "rating": 5, "price": 123.45}, {"product id": 789, "rating": 4, "price": 234.56}], "preferences": {"categories": ["electronics", "gaming"], "price range": {"min": 50, "max": 500}, "brands": ["Apple", "Samsung", "Sony"]}}
```

Response:

OK

Retrieve the whole JSON document

```
valkey-cli -h localhost -p 6379 -3 JSON.GET user:6379 $ | jq -c .'
```

Response:

```
[  
  {  
    "id": 6379,  
    "name": "Roberto Luna-Rojas",  
    "country": "Mexico M",  
    "bio": "Tech enthusiast and early adopter who loves cutting-edge gadgets...",  
    "avatar": "data:image/svg+xml;base64,PHN2ZyB2aWV3Qm94PSIwIDAgODAgODAi...",  
    "purchase_history": [  
      {"product_id": 456, "date": "2024-12-15", "rating": 5, "price": 123.45},  
      {"product_id": 789, "date": "2024-11-20", "rating": 4, "price": 234.56}  
    ],  
    "preferences": {  
      "categories": ["electronics", "gaming" ],  
      "price_range": { "min": 50, "max": 500 },  
      "brands": ["Apple", "Samsung", "Sony"]  
    }  
  }]
```

What if I want to only find products over \$100 and bellow \$200? Let's use [JSONPath](#)

```
valkey-cli -h localhost -p 6379 -3 \
JSON.GET user:6379 \
'$.purchase_history[?(@.price > 100 && @.price < 200)]' \
| jq -c '.'
```

Response:

```
[  
  {  
    "product_id": 456,  
    "date": "2024-12-15",  
    "rating": 5,  
    "price": 123.45  
  }  
]
```

If I want to update the rating for product 789, I can do so:

```
JSON.SET user:6379 \
$.purchase_history[?(@.product_id==789)].rating 4.5
```

Response:

```
OK
```

Verify the change by getting the product details:

```
valkey-cli -h localhost -p 6379 -3 \
JSON.GET user:6379 \
'$._purchase_history[?(@.product_id==789)]' \
| jq -C '.'
```

Response:

```
[  
  {  
    "product_id": 789,  
    "date": "2024-11-20",  
    "rating": 4.5,  
    "price": 234.56  
  }  
]
```

User Embedding Generation

User embeddings are created by combining:

- Bio text semantic analysis
- Purchase history patterns
- Preference indicators
- Behavioral signals

This creates a vector representation of user preferences for personalized search.

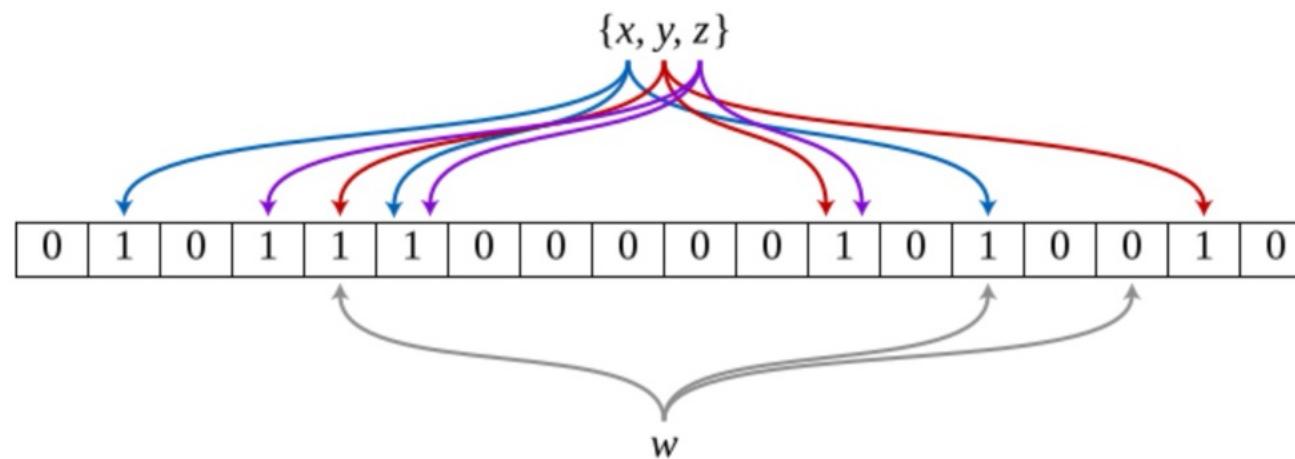
Component 3: Viewed Products Tracking with Valkey-Bloom

The Challenge

Efficiently tracking which products each user has viewed without storing massive sets that consume memory and slow down queries.

The Solution: Bloom Filters

Probabilistic data structure that provides memory-efficient membership testing.



Bloom Filters example

```
# Initialize Bloom filter for each user
bloom_key = f"viewed:{user_id}"
client.bf().reserve(bloom_key, 0.01, 1000) # 1% error rate, 1000 items capacity

# Mark product as viewed
def mark_product_viewed(user_id, product_id):
    bloom_key = f"viewed:{user_id}"
    valkey_client.bf().add(bloom_key, product_id)

# Check if product was viewed
def is_product_viewed(user_id, product_id):
    bloom_key = f"viewed:{user_id}"
    return valkey_client.bf().exists(bloom_key, product_id)

# UI Integration - show :: emoji for viewed products
for product in products:
    product['viewed'] = is_product_viewed(user_id, product['id'])
```

Why Bloom Filters?

- **Memory Efficient:** Uses minimal memory regardless of item count
- **Fast Operations:** $O(1)$ add and lookup operations
- **Scalable:** Handles millions of items with consistent performance
- **Probabilistic:** No false negatives, controlled false positive rate
- **Perfect for UX:** "Have I seen this before?" is ideal for Bloom filters

Real-World Benefits

- **Reduced Cognitive Load:** Users can quickly identify new vs. seen products
- **Improved Engagement:** Focus attention on unseen items
- **Memory Savings:** 99% less memory than storing actual sets
- **Performance:** No impact on search speed

Component 4: Vector Similarity Search with Valkey-Search

The Challenge

Finding products similar to user preferences and enabling semantic search beyond keyword matching.

The Solution: Vector Search with HNSW

Valkey-Search provides high-performance vector similarity using Hierarchical Navigable Small World (HNSW) algorithm:

Valkey-Search index creation example:

```
# Index creation with vector field
FT.CREATE products ON HASH PREFIX 1 product: SCHEMA
brand_tags TAG SEPARATOR ,
search_tags TAG SEPARATOR ,
region TAG
price NUMERIC
rating NUMERIC
embedding VECTOR HNSW 6 TYPE FLOAT32 DIM 1024 DISTANCE_METRIC COSINE
```

Hybrid search

```
# Hybrid search query combining filters + vector similarity
def search_products(user_embedding, tags, region=None):
    # Build tag filter
    tag_filter = " ".join(f"@search_tags:{{{tag}}}" for tag in tags)
    # Add region filter if specified
    if region:
        tag_filter += f" @region:{{{region}}}"
    # Combine with vector search
    query = f"({tag_filter})=>[KNN 25 @embedding $user_vec]"
    return valkey_client.ft("products").search(
        Query(query).return_fields("id", "name", "price", "rating")
            .sort_by("_score", asc=False)
            .dialect(2),
        query_params={"user_vec": user_embedding}
    )
```

Vector Similarity Concepts

What is a Vector?



A [vector](#) is a quantity that has both magnitude (size) and direction. It's a fundamental concept in mathematics and physics, used to describe quantities that can't be fully represented by a single number alone.

Vector Similarity Concepts (cont)

1. Embedding Spaces

- Products and users exist in high-dimensional vector space
- Similar items cluster together
- Distance metrics measure similarity

2. Cosine Similarity

- Measures angle between vectors, not magnitude
- Perfect for semantic similarity (0 = identical, 1 = opposite)
- Robust to vector normalization

Vector Similarity Concepts (cont)

3. HNSW Algorithm

- Hierarchical graph structure for approximate nearest neighbor search
- Logarithmic search complexity: $O(\log N)$
- Tunable accuracy vs. speed tradeoffs

4. Hybrid Search Benefits

- **Precision:** Keyword filters ensure relevance
- **Discovery:** Vector search finds unexpected matches
- **Personalization:** User embedding biases results toward preferences

Maximal Marginal Relevance (MMR)

To avoid showing too many similar products, we use MMR for result diversification: Balance relevance VS Diversity 🐍

```
def mmr_rerank(query_embedding, candidate_embeddings, lambda_param=0.7, top_n=5):
    """
    lambda_param: 1.0 = pure relevance, 0.0 = pure diversity
    """
    selected_indices = []
    # Start with most relevant item
    relevance_scores = cosine_similarity(candidate_embeddings, query_embedding)
    selected_indices.append(np.argmax(relevance_scores))
    # Iteratively select items that are relevant but diverse
    while len(selected_indices) < top_n:
        mmr_scores = {}
        for candidate_idx in remaining_candidates:
            relevance = relevance_scores[candidate_idx]
            # Measure similarity to already selected items
            diversity = max_similarity_to_selected(candidate_idx, selected_indices)
            # Balance relevance and diversity
            mmr_scores[candidate_idx] = lambda_param * relevance - (1 - lambda_param) * diversity
        # Select item with highest MMR score
        best_candidate = max(mmr_scores, key=mmr_scores.get)
        selected_indices.append(best_candidate)
    return selected_indices
```

Component 5: Personalized Recommendations

User Persona-Based Results

The system uses detailed user personas to tailor search results 

```
# Example personas
personas = {
    "101": {
        "name": "Roberto Luna-Rojas",
        "bio": "Tech enthusiast and geek by nature..",
        "interests": ["technology", "innovation", "smart_home", "gadgets"]
    },
    "102": {
        "name": "Tay Tay",
        "bio": "Best female pop artist ever...",
        "interests": ["music", "cats"]
    }
}
```

Personalization Pipeline

- 1. User Embedding:** Convert bio and preferences to vector
- 2. Product Matching:** Find products with similar embeddings
- 3. Context Filtering:** Apply user's category and price preferences
- 4. Relevance Scoring:** Combine similarity + user history + ratings
- 5. Diversification:** Use MMR to avoid redundant recommendations

Component 6: Session Management with Valkey

The Challenge

Managing user sessions, shopping carts, and temporary state across requests while maintaining performance.

The Solution: Valkey Strings and Hashes

```
# Session storage  
session_key = f"session:{session_id}"
```

```
# Store session data as hash
HSET session:abc123
    user_id 101
    cart_total 299.99
    last_activity 1704067200
    preferences '{"theme":"dark","language":"en"}'

# Shopping cart as list
LPUSH cart:abc123 "product:456" "product:789"

# Recent activity tracking
LPUSH activity:abc123 "viewed:product:456" "added_to_cart:product:789"
LTRIM activity:abc123 0 99 # Keep last 100 activities

# Session expiration
EXPIRE session:abc123 3600 # 1 hour TTL
```

Session Benefits

- **Fast Access:** O(1) session retrieval
- **Automatic Cleanup:** TTL-based session expiration
- **Atomic Updates:** Update cart without race conditions
- **Scalability:** Shared sessions across multiple app instances

Component 7: LLM Response Caching

The Challenge

AI-generated personalized product descriptions are expensive to compute and can have high latency, especially when using cloud APIs.

The Solution: Intelligent Caching Strategy

```
# Cache key structure
cache_key = f"llm_cache:user:{user_id}:product:{product_id}"
```

```
# Cache lookup before AI generation
def get_personalized_description(user_profile, product):
    cache_key = f"llm_cache:user:{user_profile['id']}:{product['id']}"
    # Try cache first
    cached_desc = valkey_client.get(cache_key)
    if cached_desc:
        return cached_desc.decode()
    # Generate new description
    prompt = f"""
        You are a helpful sales assistant. A user named {user_profile['name']}
        is considering the product: '{product['name']}'.
        Their bio is: '{user_profile['bio']}'.
        Write a personalized paragraph that addresses their interests.
        """
    # Call AI backend (AWS Bedrock, Google Gemini, or local Ollama)
    description = generate_with_ai(prompt)
    # Cache for 2 hours
    valkey_client.set(cache_key, description, ex=7200)
    return description
```

Caching Strategy Benefits

- **Cost Reduction:** Avoid repeated expensive AI API calls
- **Latency Improvement:** Cached responses return in <1ms vs 500-2000ms for AI generation
- **Scalability:** Handle more concurrent users with same AI quota
- **Reliability:** Cached responses available even if AI service is down

Cache Performance Metrics

```
# Cache hit rate monitoring
def track_cache_performance():
    total_requests = valkey_client.get("cache:total_requests") or 0
    cache_hits = valkey_client.get("cache:hits") or 0
    hit_rate = (cache_hits / total_requests) * 100 if total_requests > 0 else 0
    print(f"Cache Hit Rate: {hit_rate:.1f}%")
    print(f"Total Requests: {total_requests}")
    print(f"Cache Hits: {cache_hits}")
```

Asynchronous Cache Warming

```
def warm_cache_async(user_profile, products):
    """
    Background thread generates descriptions for products
    without blocking the user interface
    """
    def background_task():
        for product in products:
            cache_key = f"llm_cache:user:{user_profile['id']}:product:{product['id']}"
            if not valkey_client.exists(cache_key):
                # Generate and cache description
                description = generate_personalized_description(user_profile, product)
                valkey_client.set(cache_key, description, ex=7200)
    # Start background thread
    threading.Thread(target=background_task).start()
```

Performance Characteristics

Valkey-Bundle Performance Profile

Operation	Data Structure	Complexity	Typical Latency
Product Lookup	Hash	O(1)	<1ms
User Profile	JSON	O(1)	<1ms
Viewed Check	Bloom Filter	O(1)	<0.1ms
Vector Search	HNSW Index	O(log N)	2.5-10ms
Cache Lookup	String	O(1)	<0.5ms
Session Access	Hash	O(1)	<1ms

Memory Efficiency

Memory usage comparison for 1M users tracking 10K products each

Traditional Set Storage:

- $1\text{M users} \times 10\text{K products} \times 8 \text{ bytes} = 80\text{GB}$

Bloom Filter Storage:

- $1\text{M users} \times 1\text{KB per filter} = 1\text{GB}$ (98.75% memory reduction!)

False positive rate: 1% (configurable)

False negative rate: 0% (guaranteed)

Scalability Patterns

Horizontal Scaling with Valkey Cluster

```
# Cluster configuration
startup_nodes = [
    ClusterNode(host="valkey-node-1", port=6379),
    ClusterNode(host="valkey-node-2", port=6380),
    ClusterNode(host="valkey-node-3", port=6381)
]
client = ValkeyCluster(startup_nodes=startup_nodes)
# Data automatically sharded across nodes
# Hash tags ensure related data stays together
HSET {user:101}:profile name "Roberto Luna-Rojas"
HSET {user:101}:session cart_total 299.99
BF.ADD {user:101}:viewed product:456
```

Development and Deployment

Local Development Setup

```
# Start Valkey-bundle
docker run -d --rm --name valkey-demo -p 6379:6379 valkey/valkey-bundle

# Setup Python environment
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt

# Initialize data
python3 load_data.py
python3 init_bloom_filters.py

# Run application
flask run --host=0.0.0.0 --port=5001
```

Production Considerations

- **Memory Planning:** Size Valkey instances based on dataset and cache requirements
- **Backup Strategy:** Regular RDB snapshots + AOF for durability
- **Monitoring:** Track cache hit rates, search latency, and memory usage
- **Security:** Network isolation, authentication, and encryption in transit

Key Takeaways

Why Valkey-Bundle Excels for Modern Applications

- 1. Unified Platform:** Single solution for diverse data needs
- 2. Performance:** Sub-millisecond operations for most use cases
- 3. Scalability:** Horizontal scaling with cluster mode
- 4. Flexibility:** Multiple data structures for different patterns
- 5. AI Integration:** Native vector search for ML applications
- 6. Developer Experience:** Rich ecosystem and tooling

When to Choose Valkey-Bundle

Perfect For:

- Real-time applications requiring low latency
- AI/ML applications with vector similarity search
- Applications with diverse data access patterns
- High-performance caching layers
- Session management and user state
- Analytics and recommendation engines

When to Choose Valkey-Bundle (cont)

⚠ Consider Alternatives For:

- Applications requiring strong consistency guarantees
- Complex relational queries with joins
- Long-term analytical data warehousing
- Applications with minimal performance requirements

The Future of Low-Latency Applications

Valkey-bundle represents the evolution toward:

- **Unified Data Platforms:** Reducing operational complexity
- **AI-Native Infrastructure:** Built-in support for vector operations
- **Edge Computing:** Fast, local data processing
- **Real-Time Personalization:** Instant, context-aware experiences

Conclusion

This demonstration showcases how Valkey-bundle's integrated approach solves real-world challenges in modern application development. By combining traditional data structures with advanced capabilities like vector search and probabilistic filters, developers can build sophisticated, high-performance applications with a single, unified platform.

The key insight is that modern applications require diverse data access patterns - from simple key-value lookups to complex vector similarity searches. Valkey-bundle provides all these capabilities in a cohesive, high-performance package that scales from prototype to production.

Conclusion (cont)

Whether you're building recommendation engines, real-time analytics, or AI-powered applications, Valkey-bundle offers the performance, flexibility, and developer experience needed for success in today's demanding application landscape.

For more information about Valkey-bundle and to explore the complete source code of this demonstration, visit: <https://valkey.io/blog/valkey-bundle-one-stop-shop-for-low-latency-modern-applications/> by Roberto Luna-Rojas

Demo enhanced from original [Valkey Search Demo](#) by Ping Xie [PingXie](#)